



University of
Pittsburgh

Introduction to Operating Systems

CS 1550



Spring 2023

Sherif Khattab

ksm73@pitt.edu

(Some slides are from **Silberschatz, Galvin and Gagne ©2013**)

Announcements

- Upcoming deadlines
 - Homework 6 is due **this Friday**
 - Quiz 1 is due on Tuesday 2/28 at 11:59 pm
 - Lab 2 is due on Tuesday 2/28 at 11:59 pm
 - Project 2 is due Friday 3/17 at 11:59 pm
- Midterm exam on Thursday 3/2
 - In-person, on paper, closed book
 - Study guide, old exam, and practice Midterm on Canvas
- Lost points because autograder or simple mistake?
 - please reach out to Grader TA over Piazza
- Navigating the Panopto Videos
 - Video contents
 - Search in captions

Previous Lecture

- Review on Bounded Buffer and Readers-Writers
- Sleepy Barbers solution using Semaphores

Today's Agenda

- Sleepy Barbers solution using Condition Variables
- How to implement condition variables
- Reflections on using semaphores and condition variables
- CPU Scheduling

Solution using Mutex and Condition Variables

- <https://cs1550-2214.github.io/cs1550-code-handouts/ProcessSynchronization/Slides/>

How to implement Condition Variables?

- How to implement condition variables?
- Reflect more on all the solutions/problems that we have studied

User-level implementation of Condition Variables

A Lock with two waiting queues

```
struct Lock {  
    Semaphore mutex(1);  
    Semaphore next(0);  
    int nextCount = 0;  
}
```

```
Acquire(){
```

```
    mutex.down();  
}
```

```
Release(){
```

```
    if(nextCount > 0){  
        next.up();  
        nextCount--;  
    } else mutex.up();  
}
```

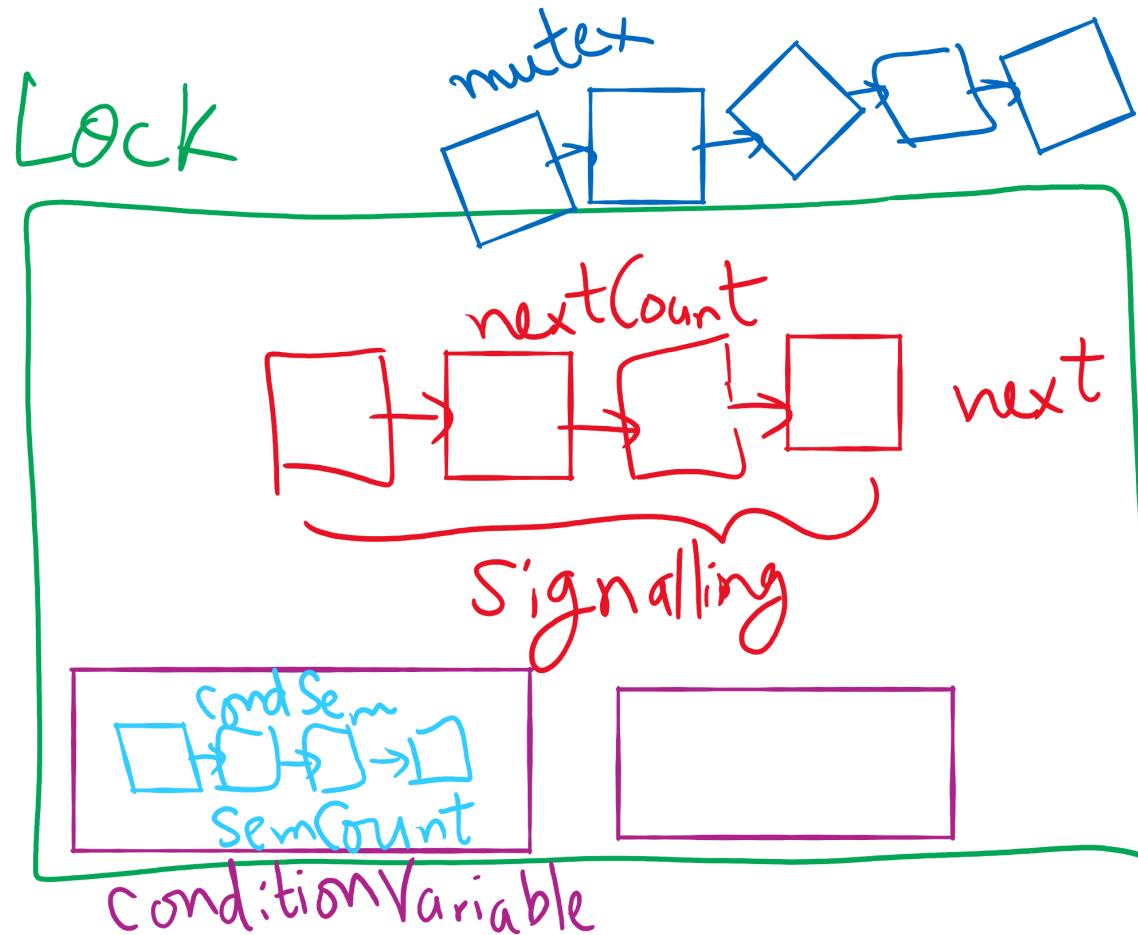
Condition Variable

```
struct ConditionVariable {  
    Semaphore condSem(0);  
    int semCount = 0;  
    Lock *lk;  
}
```

```
Wait(){  
    if(lk->nextCount > 0)  
        lk->next.up();  
        lk->nextCount--;  
  
    else {  
        lk->mutex.up();  
    }  
  
    semCount++;  
    condSem.down();  
    semCount--;  
}
```

```
Signal(){  
    if(semCount > 0){  
        condSem.up()  
        lk->nextCount++  
        lk->next.down();  
        lk->nextCount—  
    }  
}
```

Lock and Condition Variable Implementation



Implementing locks with semaphores

- Use mutex to ensure exclusion within the lock bounds
- Use next to give lock to processes with a higher priority (why?)
- nextCount indicates whether there are any higher priority waiters

```
class Lock {  
    Semaphore mutex(1);  
    Semaphore next(0);  
    int nextCount = 0;  
};
```

```
Lock::Acquire()  
{  
    mutex.down();  
}
```

```
Lock::Release()  
{  
    if (nextCount > 0)  
        next.up();  
    else  
        mutex.up();  
}
```

Implementing condition variables

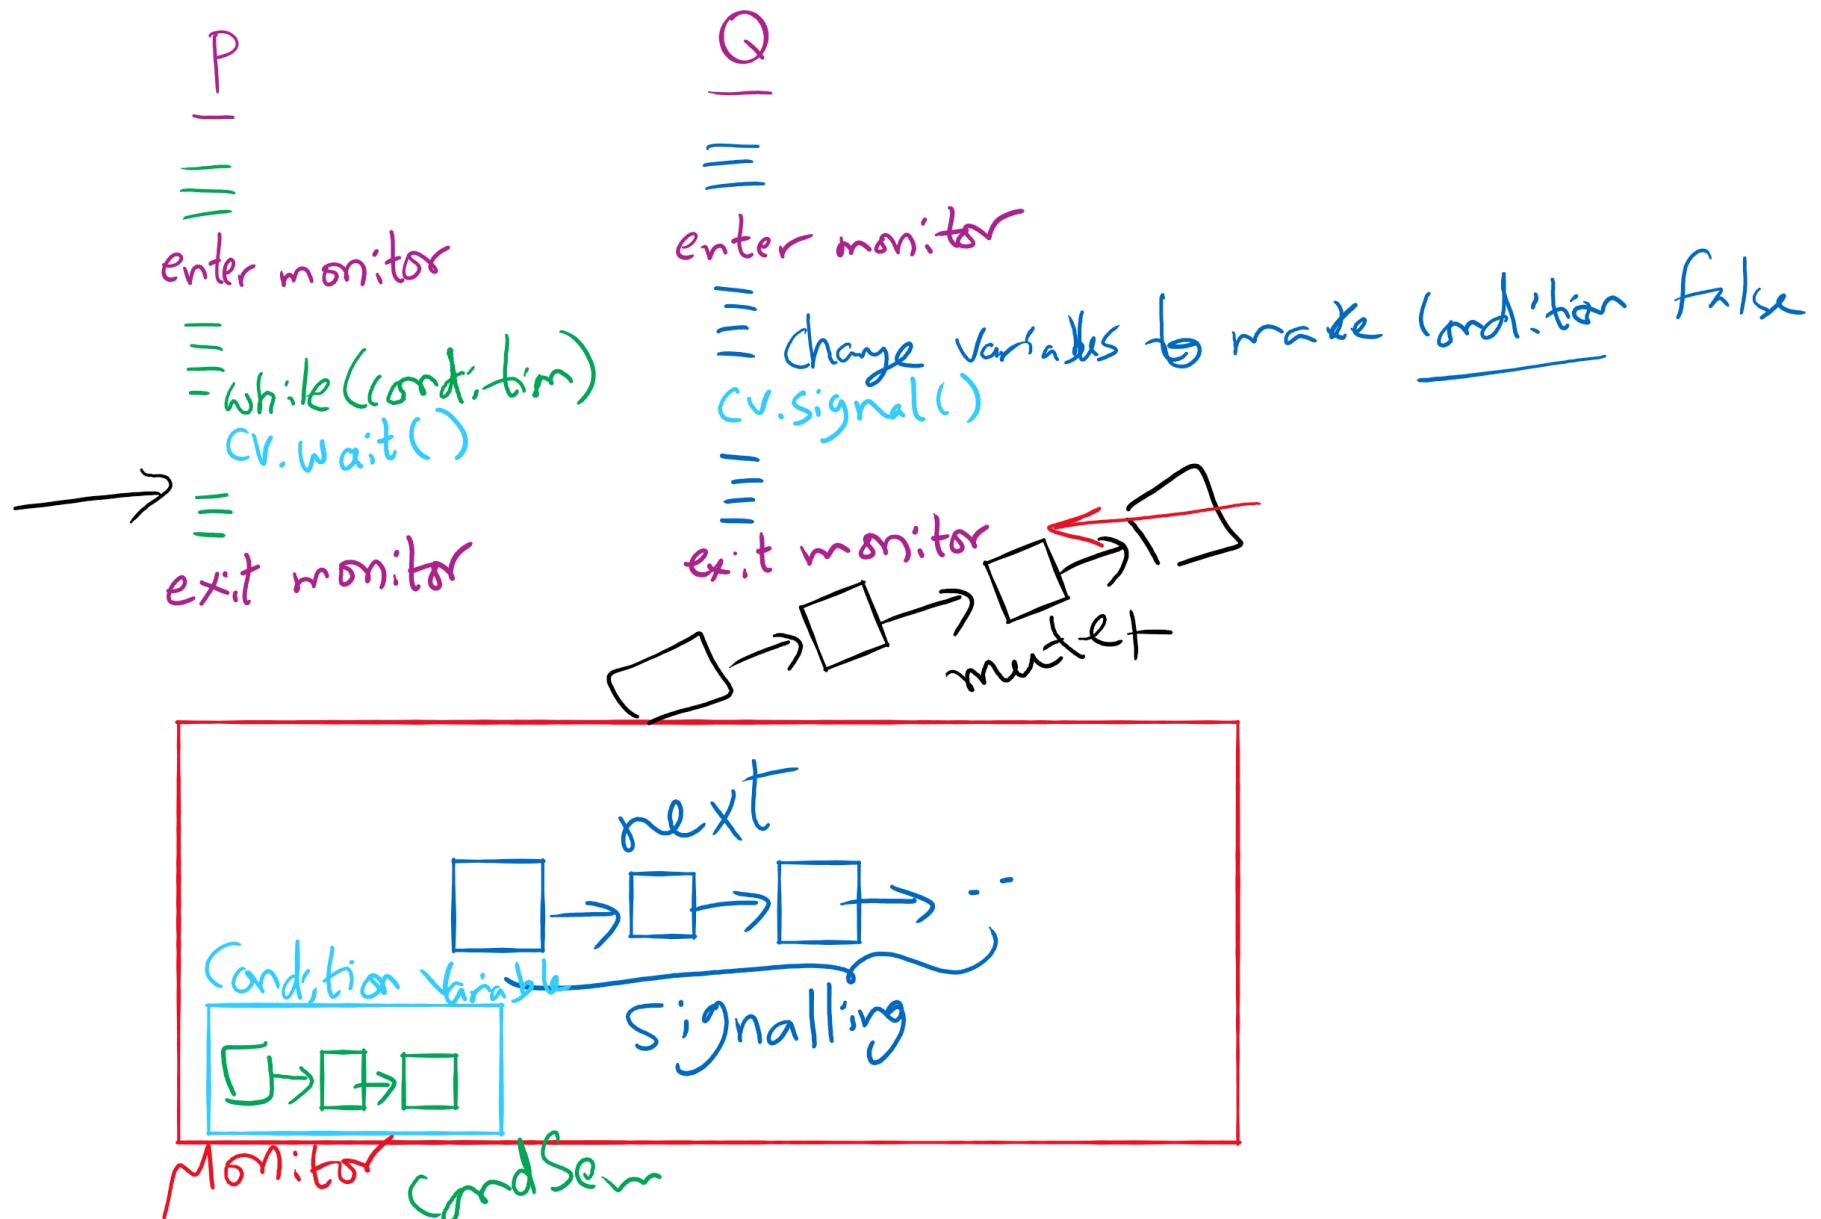
- Are these Hoare or Mesa semantics?
- Can there be multiple condition variables for a single Lock?

```
class Condition {  
    Lock *lock;  
    Semaphore condSem(0);  
    int semCount = 0;  
};
```

```
Condition::Wait ()  
{  
    semCount += 1;  
    if (lock->nextCount > 0)  
        lock->next.up();  
    else  
        lock->mutex.up();  
    condSem.down ();  
    semCount -= 1;  
}
```

```
Condition::Signal ()  
{  
    if (semCount > 0) {  
        lock->nextCount += 1;  
        condSem.up ();  
        lock->next.down ();  
        lock->nextCount -= 1;  
    }  
}
```

Process Synchronization inside Monitors



Reflections on semaphore usage

- Semaphores can be used as
 - Resource counters
 - Waiting spaces
 - For mutual exclusion

Reflections on Condition Variables

- Define a class and put all shared variables inside the class
- Include a mutex and a condition variable in the class
- For each public method of the class
 - Start by locking the mutex lock
 - If need to wait, use a while loop and wait on the condition variable
 - Before broadcasting on the condition variable, make sure to change the waiting condition

Final Remarks on Process Synchronization

- Many other synchronization mechanisms
 - Message passing
 - Barriers
 - Futex
 - Re-entrant locks
 - Atomic*

Problem of the Day: CPU Scheduling

How does the ***short-term scheduler*** select the next process to run?

CPU Scheduling

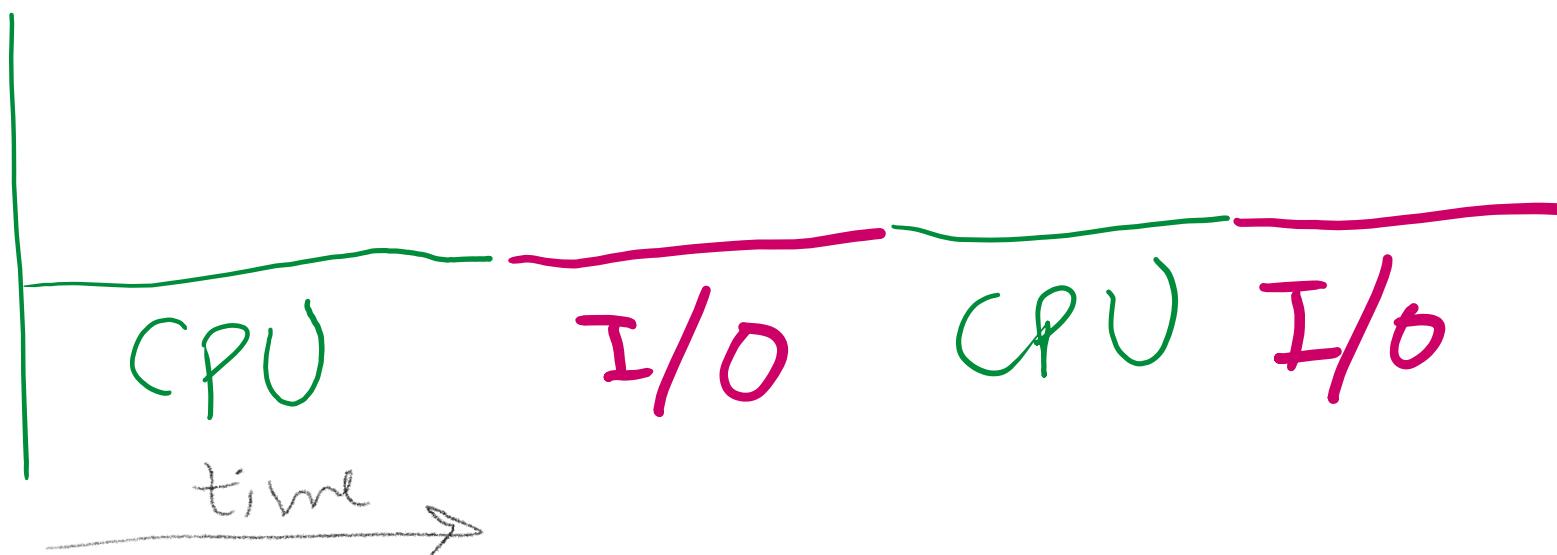
- Scheduling the processor among all ready processes
- User-oriented criteria
 - **Response Time**: Elapsed time between the submission of a request and the receipt of a response
 - **Turnaround Time**: Elapsed time between the submission of a process to its completion
- System-oriented criteria
 - Processor utilization
 - Throughput: number of process completed per unit time
 - Fairness

Short-Term Scheduler Dispatcher

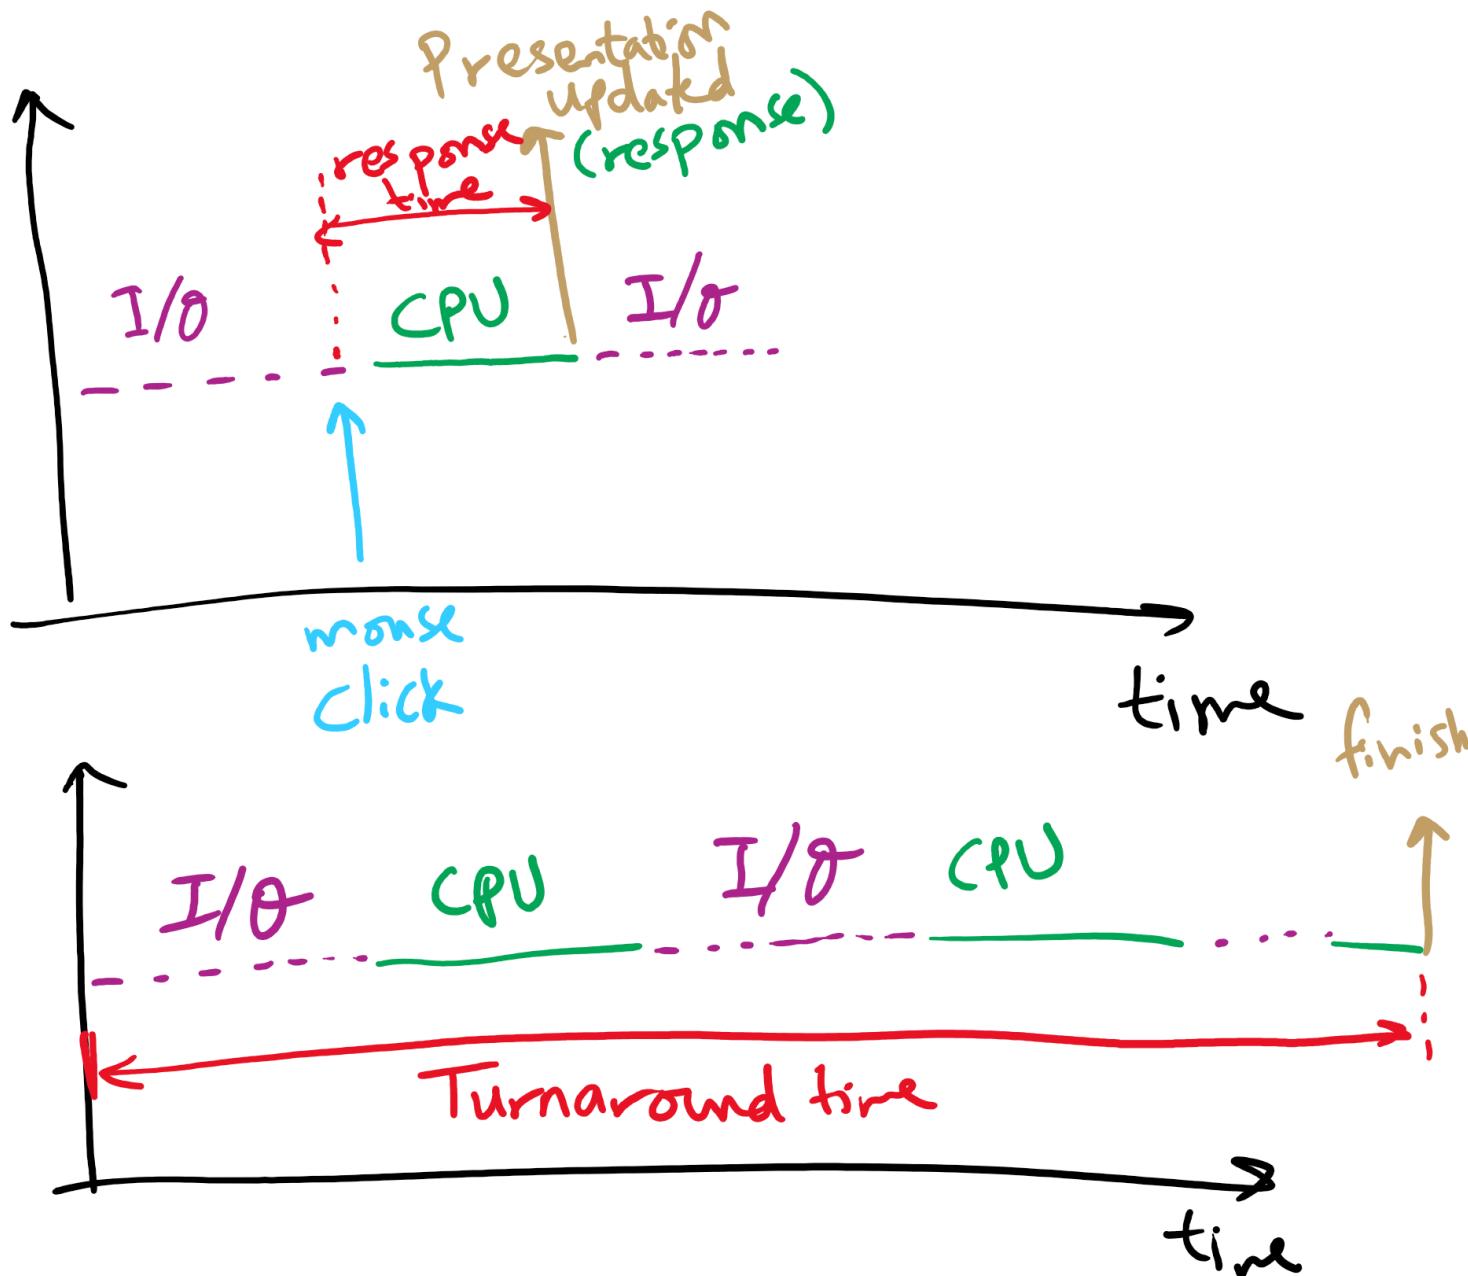
- The dispatcher is the module that gives control of the CPU to the process selected by the short-term scheduler
- The functions of the dispatcher include:
 - Switching context
 - Switching to user mode
 - Jumping to the location in the user program to restart execution
- The dispatch latency must be minimal

The CPU-I/O Cycle

- Processes require alternate use of processor and I/O in a repetitive fashion
- Each cycle consist of a CPU burst followed by an I/O burst
 - A process terminates on a CPU burst
- CPU-bound processes have longer CPU bursts than I/O-bound processes



Response time vs. Turnaround time



Scheduling Algorithms

- First-Come, First-Served Scheduling
- Shortest-Job-First Scheduling
 - Also referred to as Shortest Process Next
- Priority Scheduling
- Round-Robin Scheduling
- Multilevel Queue Scheduling
- Multilevel Feedback Queue Scheduling

Characterization of Scheduling Policies

- The **selection function** determines which ready process is selected next for execution
- The **decision mode** specifies the instants in time the selection function is exercised
 - Nonpreemptive
 - Once a process is in the running state, it will continue until it terminates or blocks for an I/O
 - Preemptive
 - Currently running process may be interrupted and moved to the Ready state by the OS
 - Prevents one process from monopolizing the processor

Process Mix Example

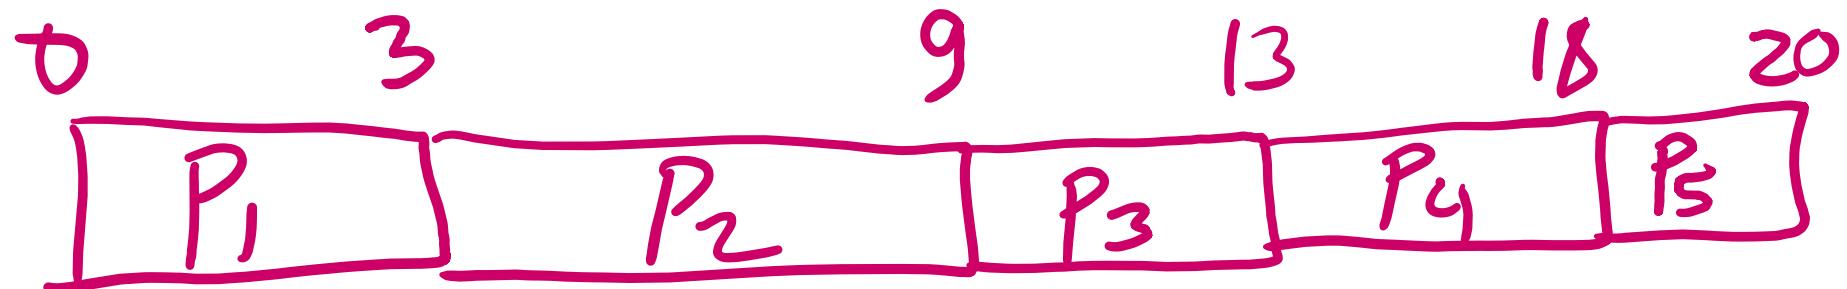
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2

Service time = total processor time needed in one (CPU-I/O) cycle
Jobs with long service time are CPU-bound jobs and are referred to as “long jobs”

First Come First Served (FCFS)

- Selection function: the process that has been waiting the longest in the ready queue (hence, FCFS)
- Decision mode: non-preemptive
 - a process runs until it blocks for an I/O

Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



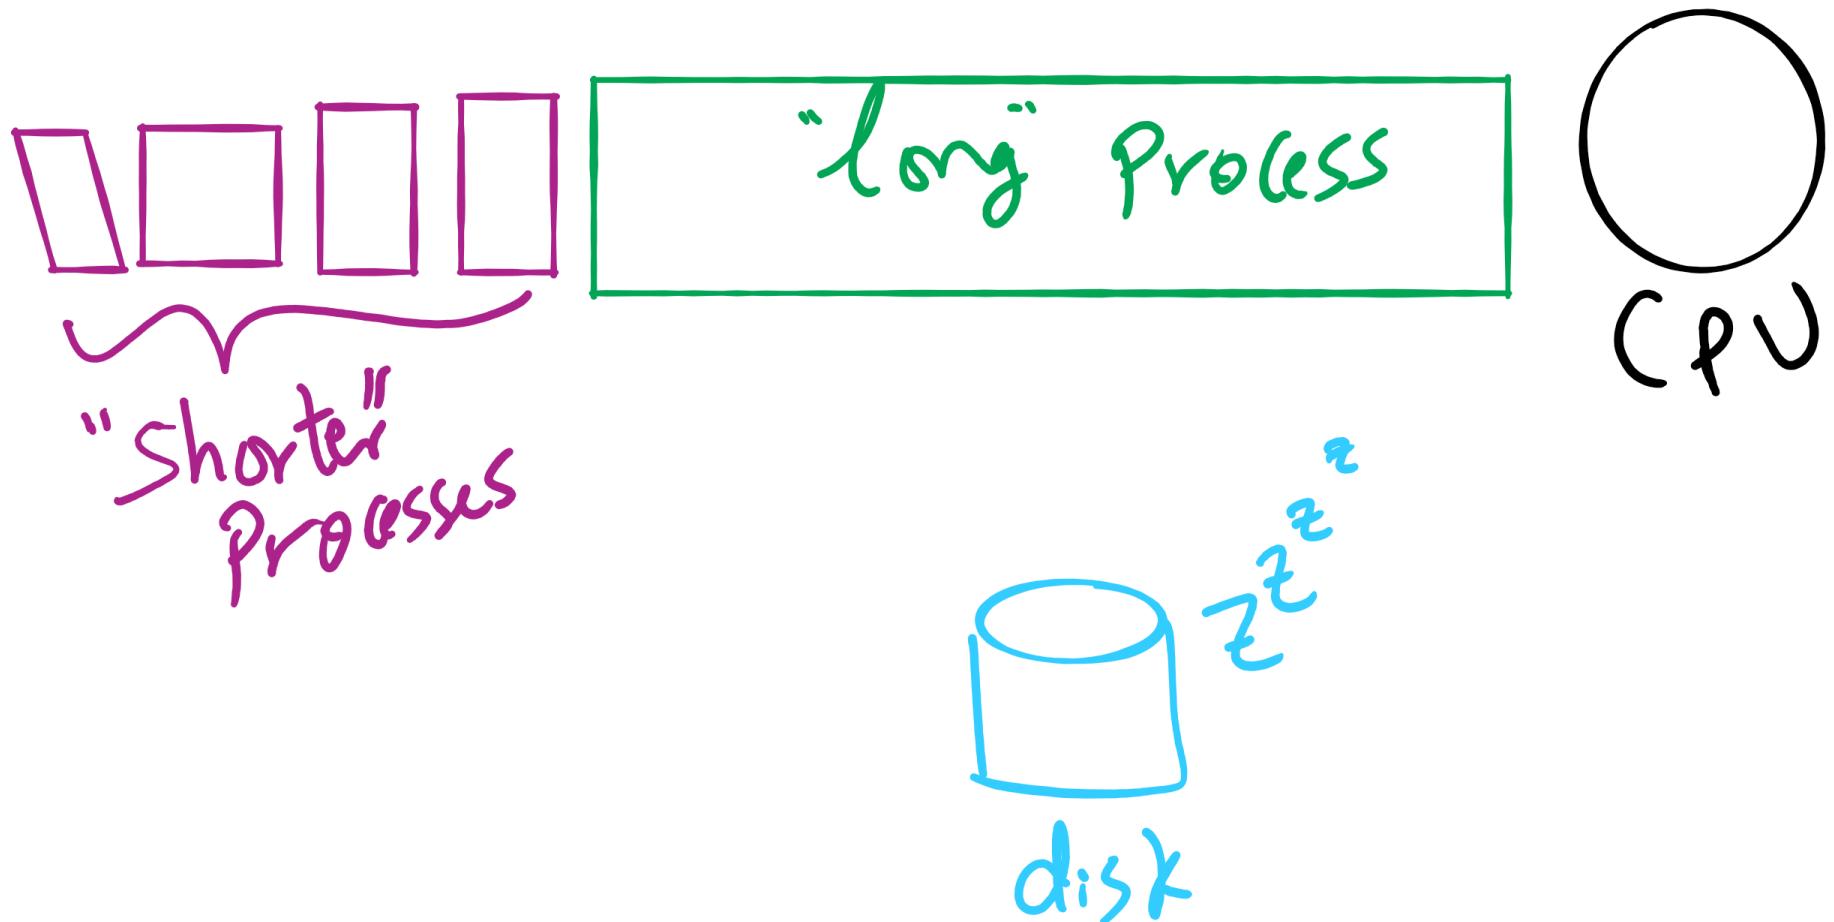
Average Response Time

$$\text{Average Response Time} = \frac{(3-0) + (9-2) + (13-4) + (18-6) + (20-8)}{5}$$

FCFS drawbacks

- Favours CPU-bound processes
 - CPU-bound processes monopolize the processor
 - I/O-bound processes have to wait until completion of CPU-bound process
 - I/O-bound processes may have to wait even after their I/Os are completed (poor device utilization)
 - Convoy effect
 - Better I/O device utilization could be achieved if I/O bound processes had higher priority

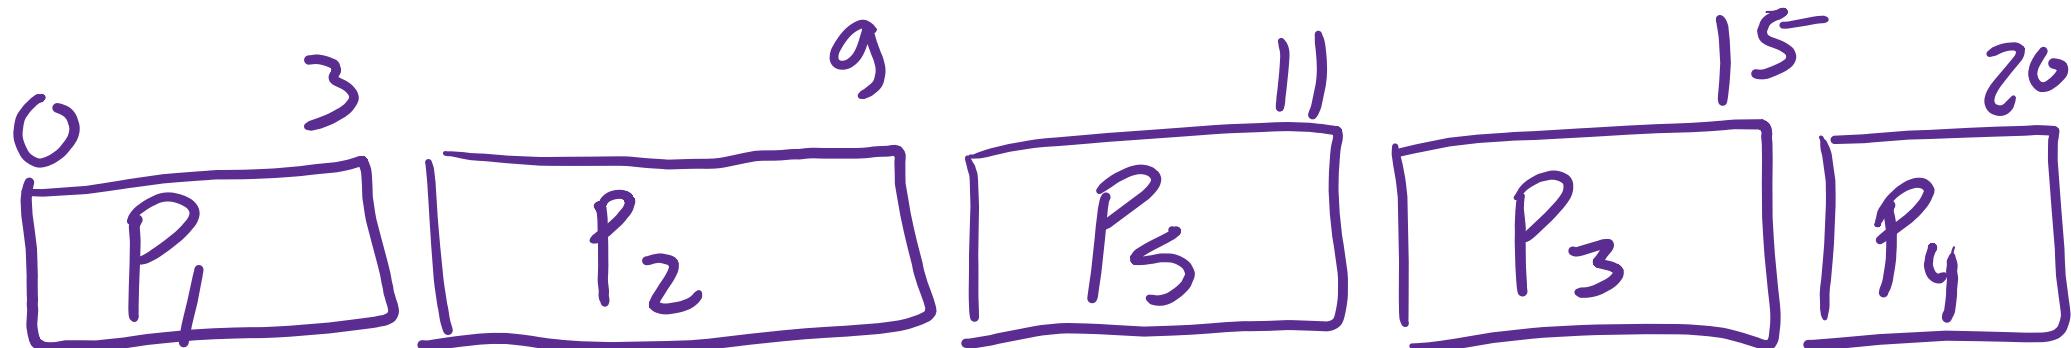
Convoy Effect



Shortest Job First (Shortest Process Next)

- Selection function: the process with the shortest expected CPU burst time
 - I/O-bound processes will be selected first
- Decision mode: non-preemptive
- The required processing time, i.e., the CPU burst time, must be estimated for each process

Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



SJF / SPN Critique

- Possibility of starvation for longer processes
- Lack of preemption is not suitable in a time sharing environment
- SJF/SPN implicitly incorporates priorities
 - Shortest jobs are given preference
 - CPU bound processes have lower priority, but a process doing no I/O could still monopolize the CPU if it is the first to enter the system

Priorities

- Implemented by having multiple ready queues to represent each level of priority
- Scheduler selects the process of a higher priority over one of lower priority
- Lower-priority may suffer starvation
- To alleviate starvation allow dynamic priorities
 - The priority of a process changes based on its age or execution history

Round-Robin

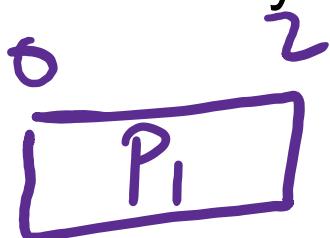
- Selection function: same as FCFS
- Decision mode: pre-emptive

- ◆ a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired
- ◆ a clock interrupt occurs and the running process is put on the ready queue

Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2

Round-Robin

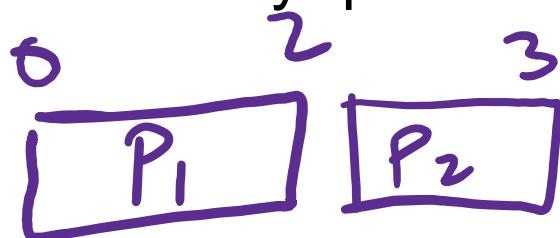
- Selection function: same as FCFS
- Decision mode: pre-emptive
 - ◆ a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired
 - ◆ a clock interrupt occurs and the running process is put on the ready queue



Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2

Round-Robin

- Selection function: same as FCFS
- Decision mode: pre-emptive
 - ◆ a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired
 - ◆ a clock interrupt occurs and the running process is put on the ready queue

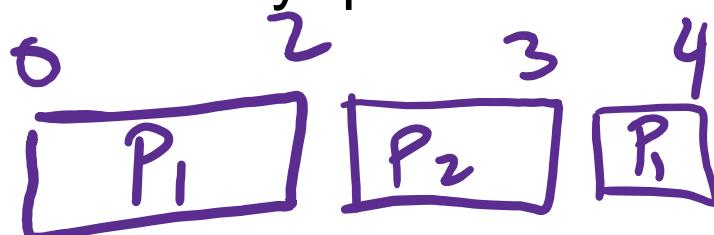


Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2

Round-Robin

- Selection function: same as FCFS
- Decision mode: pre-emptive
 - ◆ a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired
 - ◆ a clock interrupt occurs and the running process is put on the ready queue

Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



Round-Robin

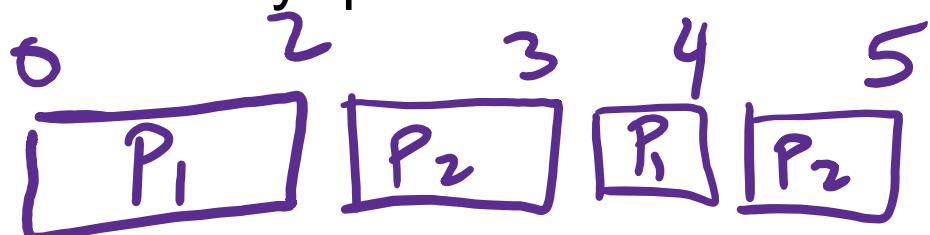
- Selection function: same as FCFS

- Decision mode: pre-emptive

- ◆ a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- ◆ a clock interrupt occurs and the running process is put on the ready queue

Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



Round-Robin

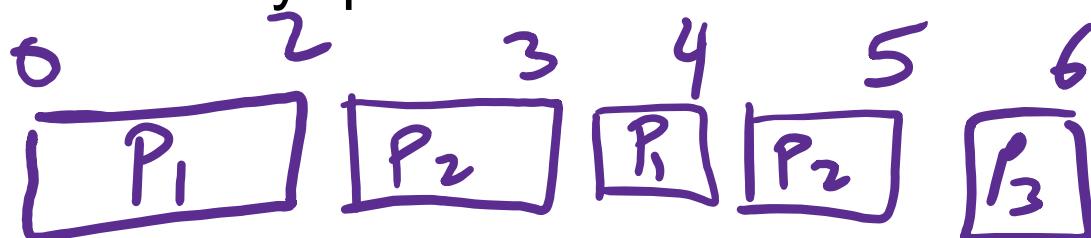
- Selection function: same as FCFS

- Decision mode: pre-emptive

- ◆ a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- ◆ a clock interrupt occurs and the running process is put on the ready queue

Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



Round-Robin

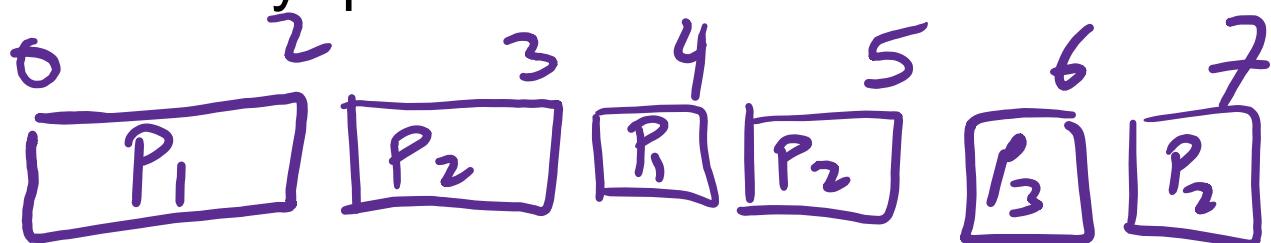
- Selection function: same as FCFS

- Decision mode: pre-emptive

- ◆ a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- ◆ a clock interrupt occurs and the running process is put on the ready queue

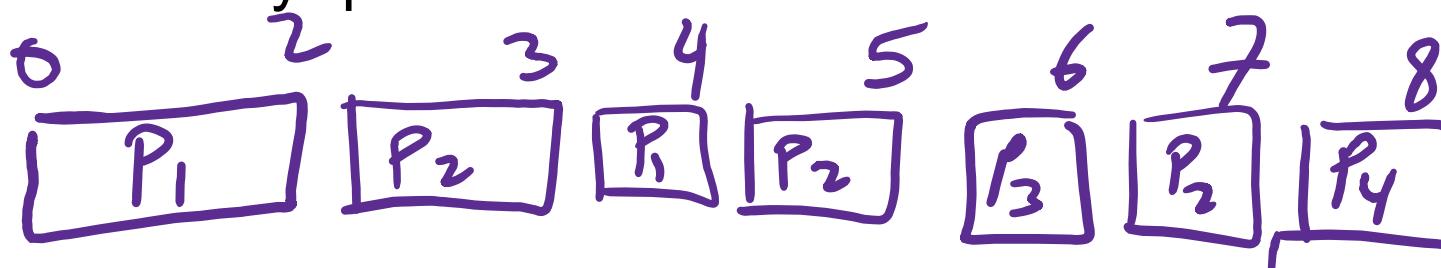
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



Round-Robin

- Selection function: same as FCFS
- Decision mode: pre-emptive
 - ◆ a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired
 - ◆ a clock interrupt occurs and the running process is put on the ready queue

Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



Round-Robin

- Selection function: same as FCFS
- Decision mode: pre-emptive

- ◆ a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- ◆ a clock interrupt occurs and the running process is put on the ready queue

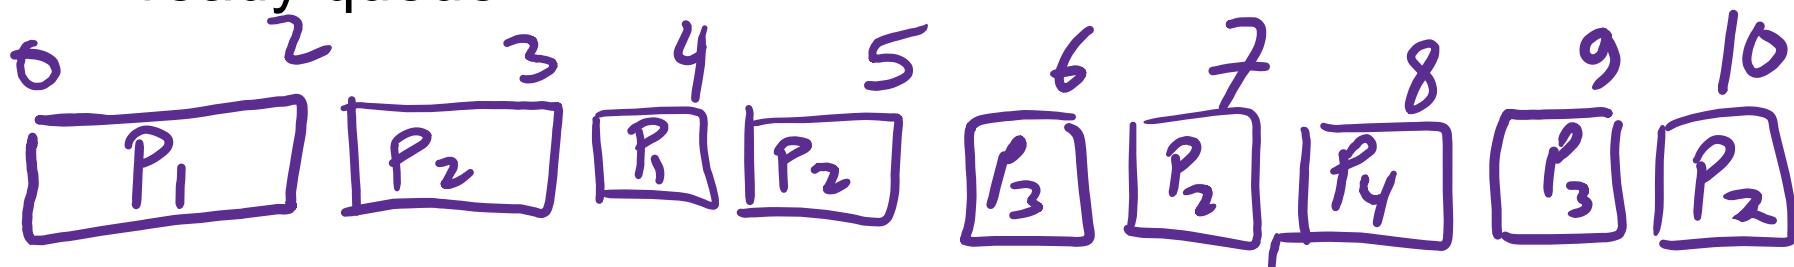
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



Round-Robin

- Selection function: same as FCFS
- Decision mode: pre-emptive
 - ◆ a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired
 - ◆ a clock interrupt occurs and the running process is put on the ready queue

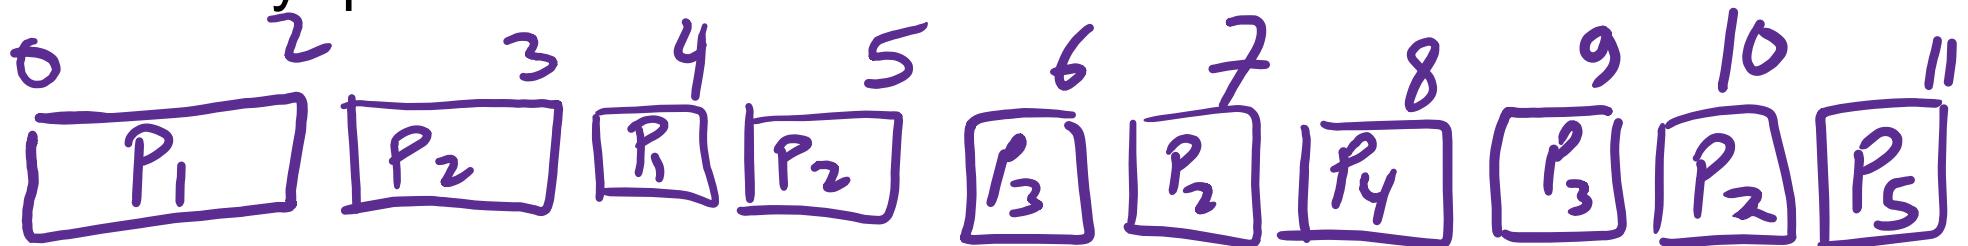
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



Round-Robin

- Selection function: same as FCFS
- Decision mode: pre-emptive
 - ◆ a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired
 - ◆ a clock interrupt occurs and the running process is put on the ready queue

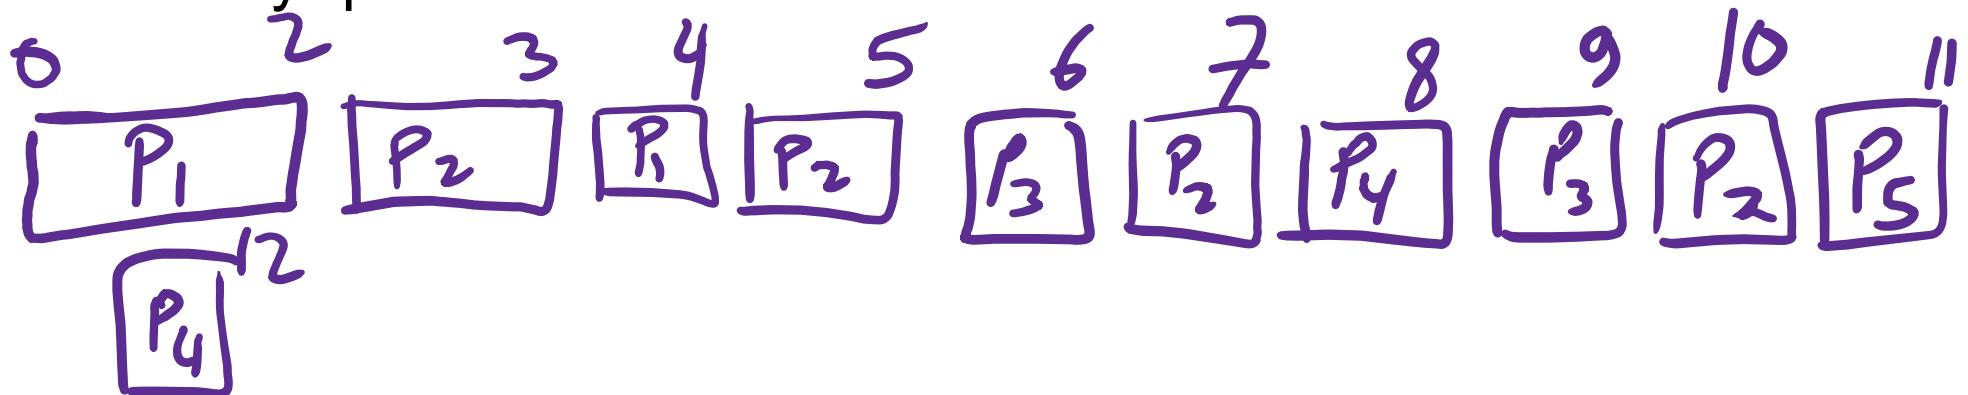
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



Round-Robin

- Selection function: same as FCFS
- Decision mode: pre-emptive
 - ◆ a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired
 - ◆ a clock interrupt occurs and the running process is put on the ready queue

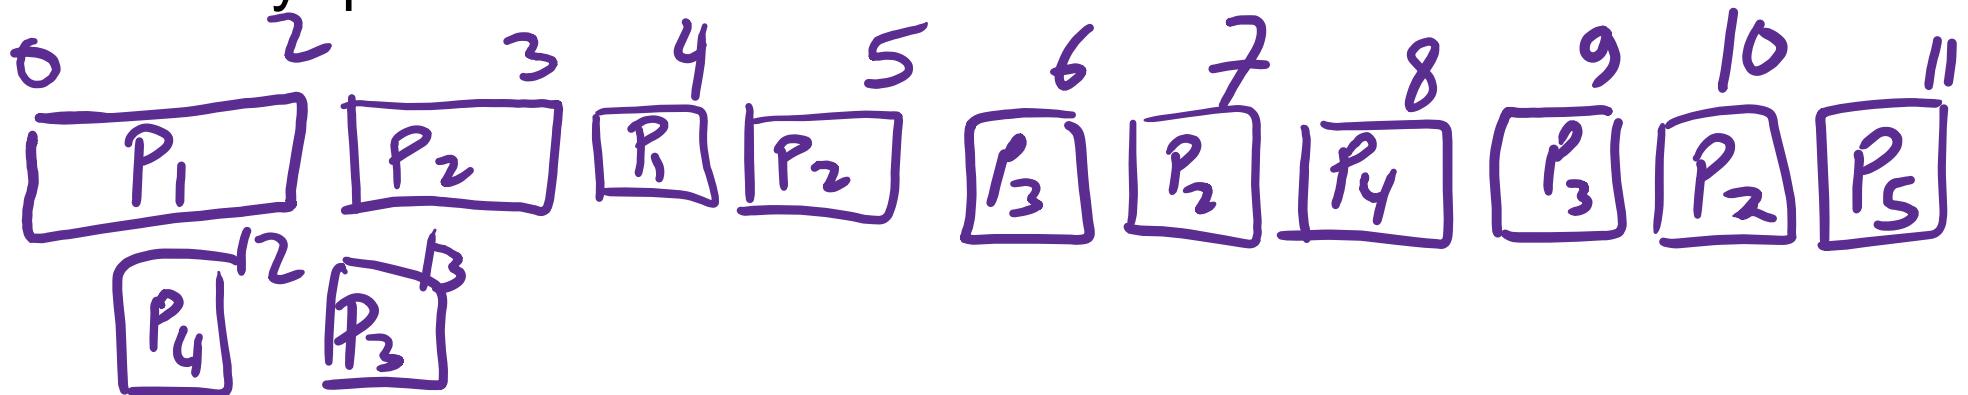
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



Round-Robin

- Selection function: same as FCFS
- Decision mode: pre-emptive
 - ◆ a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired
 - ◆ a clock interrupt occurs and the running process is put on the ready queue

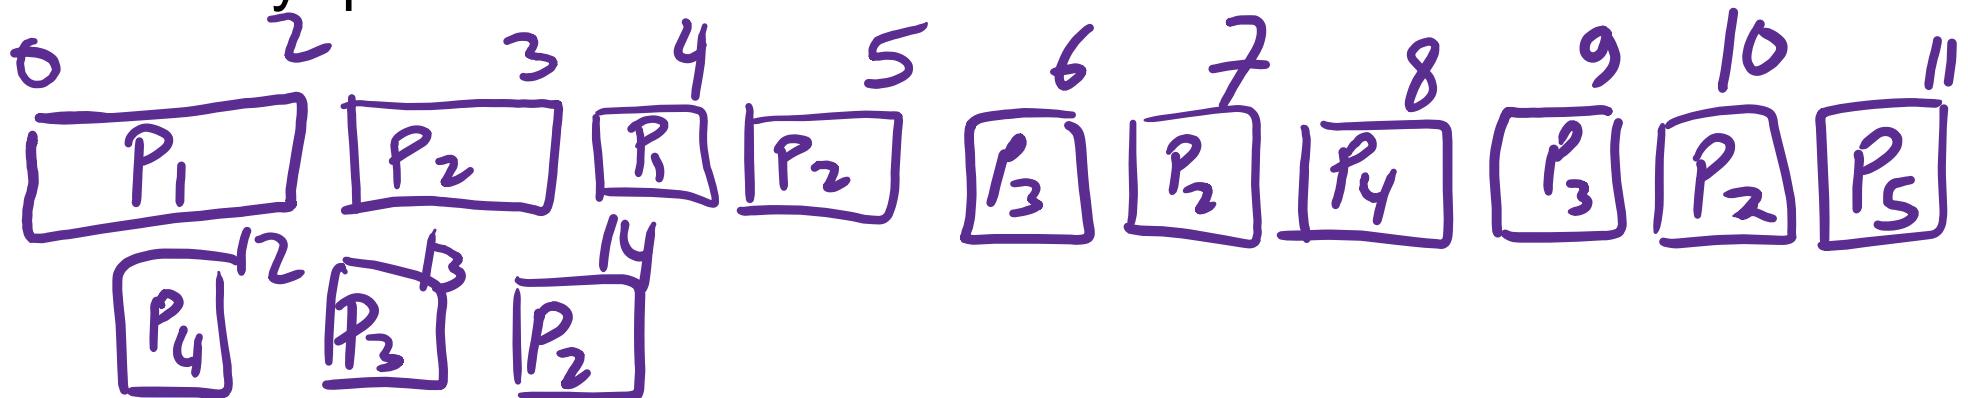
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



Round-Robin

- Selection function: same as FCFS
- Decision mode: pre-emptive
 - ◆ a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired
 - ◆ a clock interrupt occurs and the running process is put on the ready queue

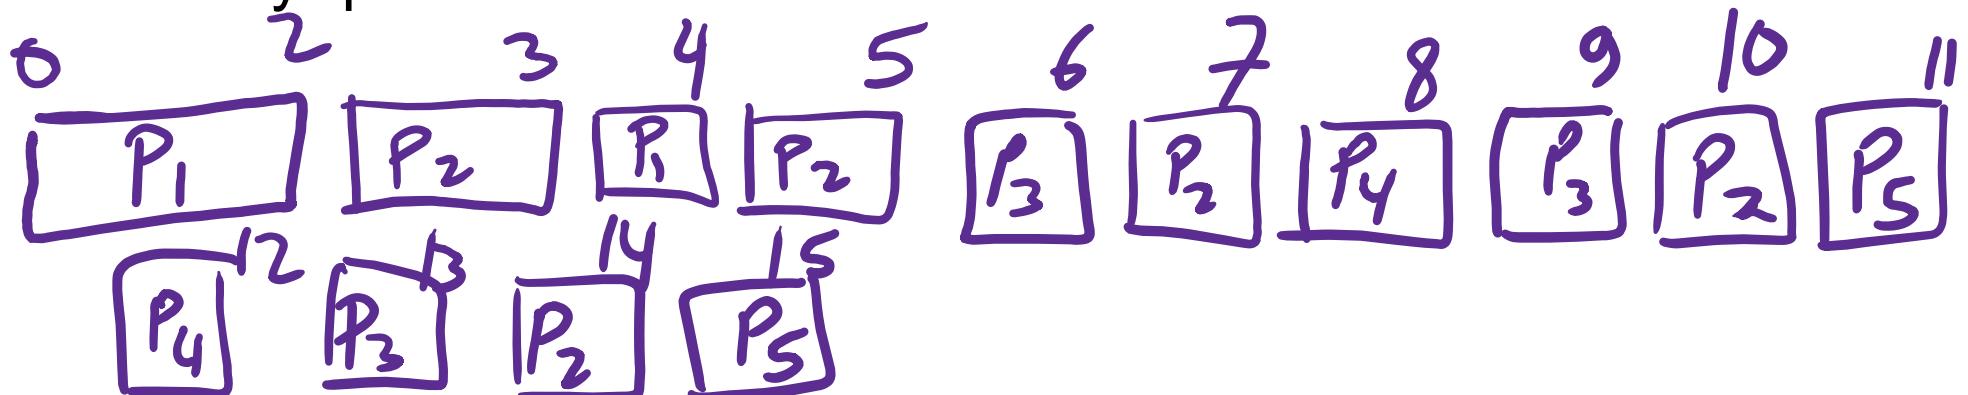
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



Round-Robin

- Selection function: same as FCFS
- Decision mode: pre-emptive
 - ◆ a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired
 - ◆ a clock interrupt occurs and the running process is put on the ready queue

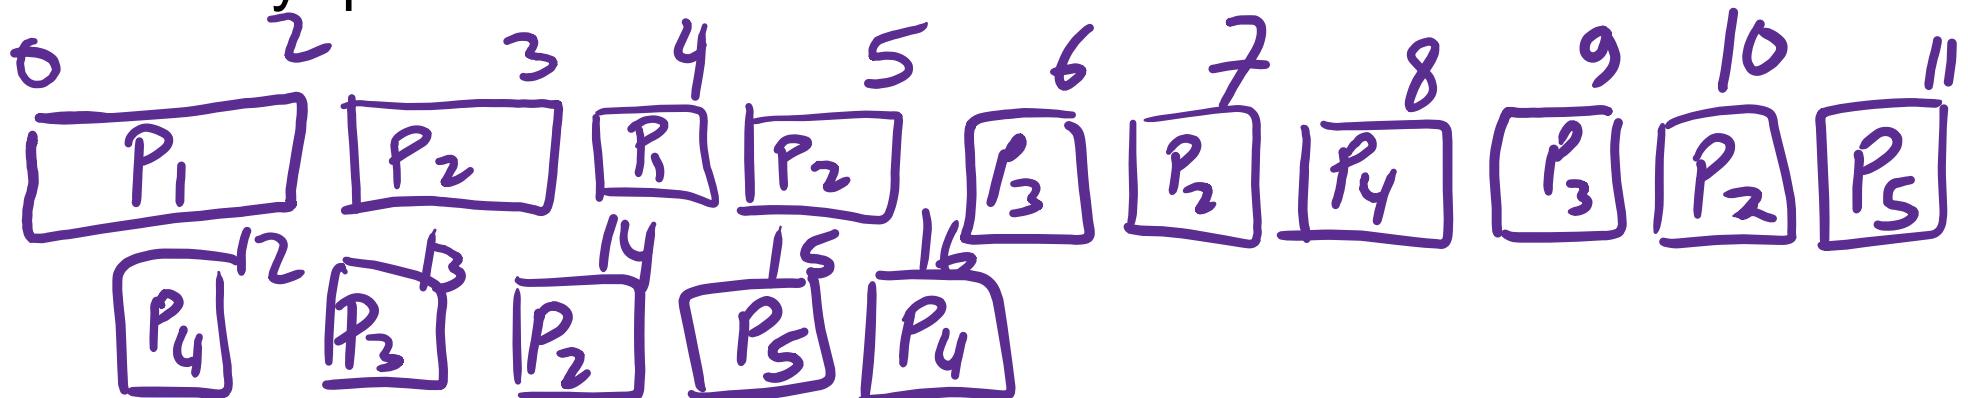
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



Round-Robin

- Selection function: same as FCFS
- Decision mode: pre-emptive
 - ◆ a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired
 - ◆ a clock interrupt occurs and the running process is put on the ready queue

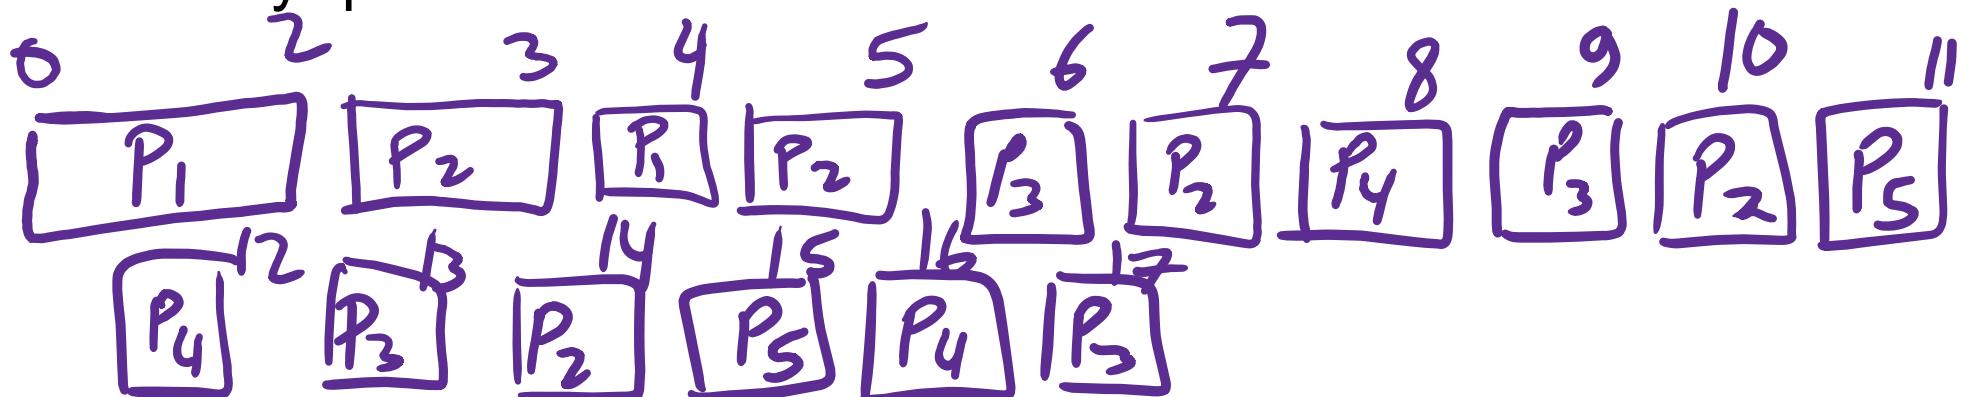
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



Round-Robin

- Selection function: same as FCFS
- Decision mode: pre-emptive
 - ◆ a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired
 - ◆ a clock interrupt occurs and the running process is put on the ready queue

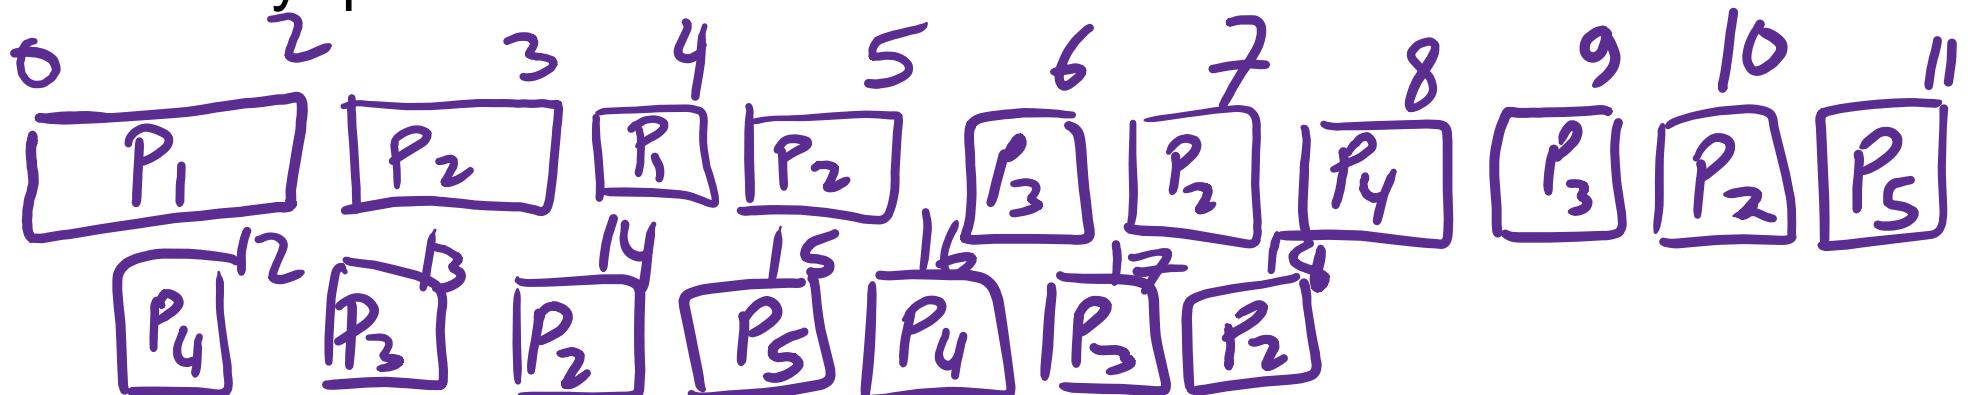
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



Round-Robin

- Selection function: same as FCFS
- Decision mode: pre-emptive
 - ◆ a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired
 - ◆ a clock interrupt occurs and the running process is put on the ready queue

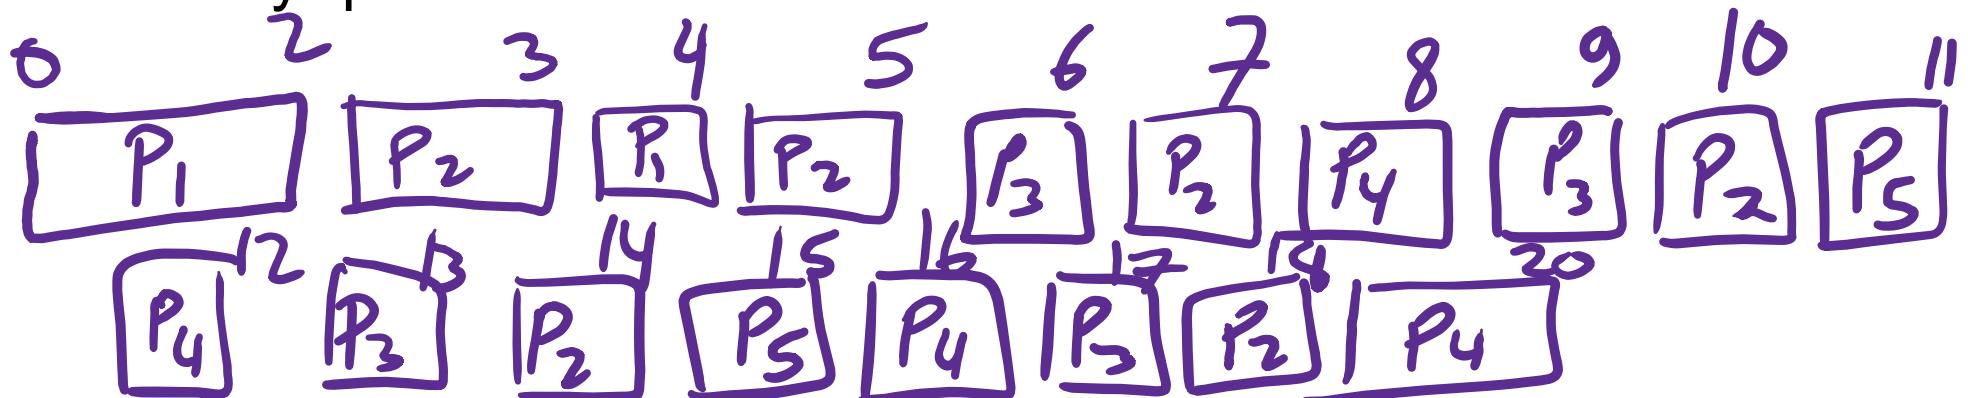
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



Round-Robin

- Selection function: same as FCFS
- Decision mode: pre-emptive
 - ◆ a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired
 - ◆ a clock interrupt occurs and the running process is put on the ready queue

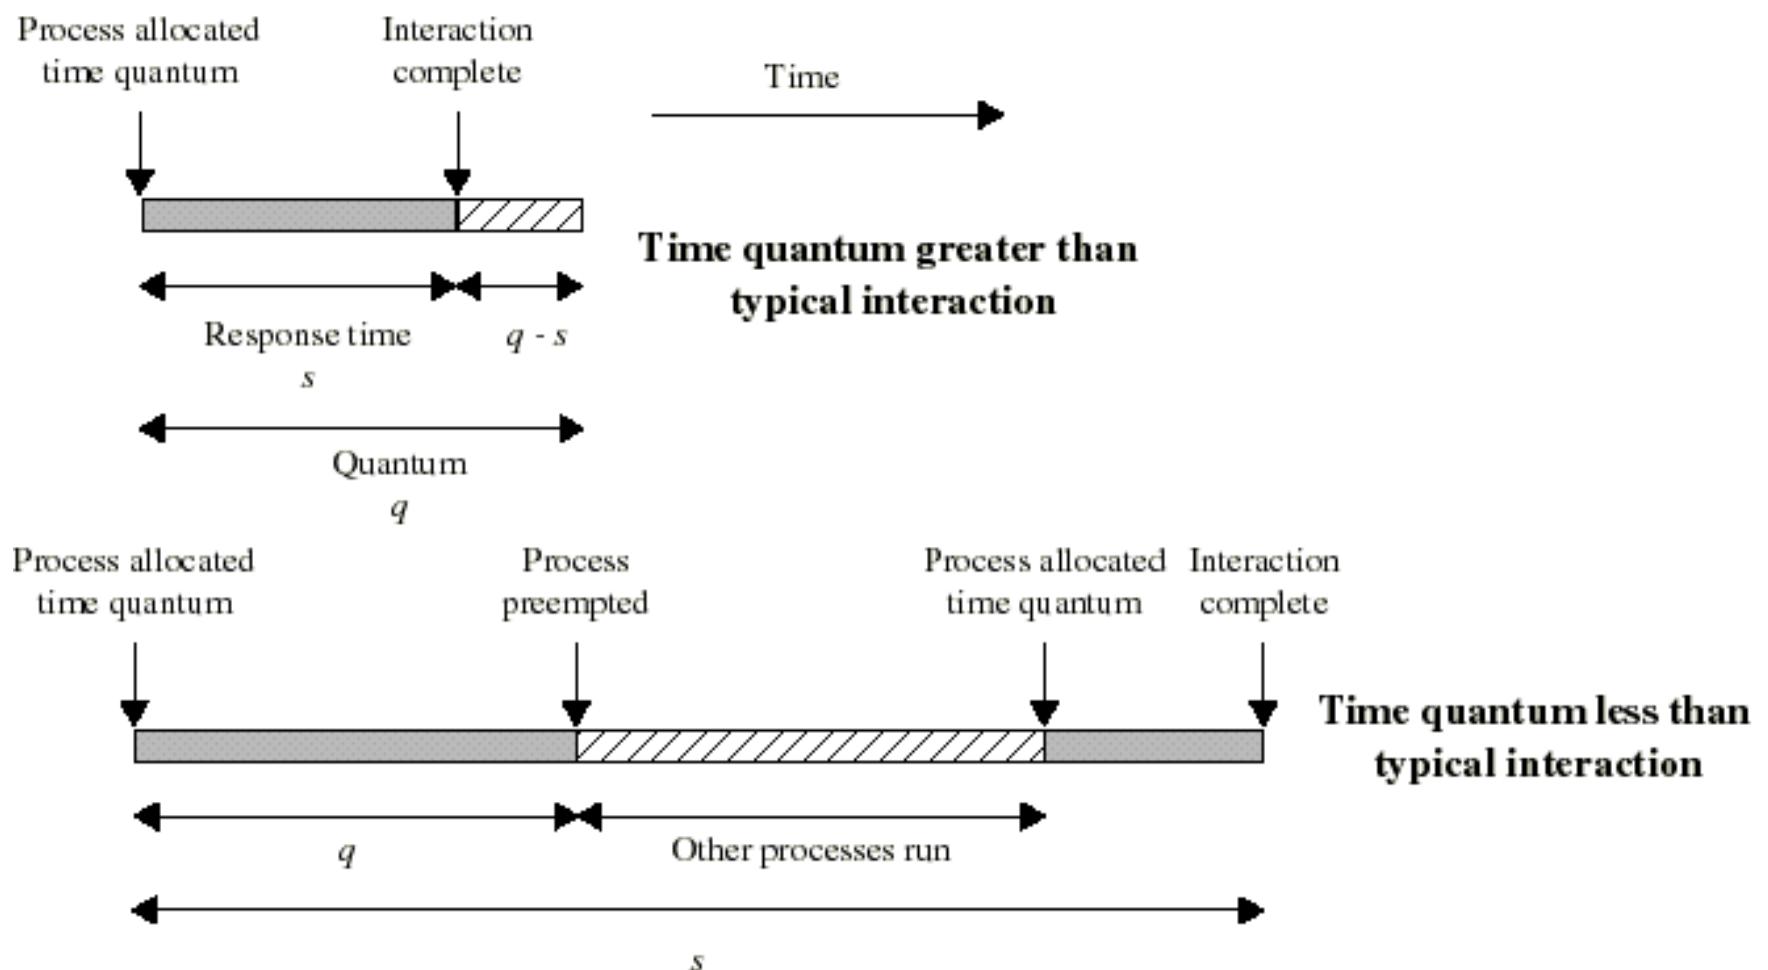
Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



RR Time Quantum

- Quantum must be substantially larger than the time required to handle the clock interrupt and dispatching
- Quantum should be larger then the typical interaction
 - but not much larger, to avoid penalizing I/O bound processes

RR Time Quantum



Quantum Length



Round Robin: critique

- Still favors CPU-bound processes
 - An I/O bound process uses the CPU for a time less than the time quantum before it is blocked waiting for an I/O
 - A CPU-bound process runs for all its time slice and is put back into the ready queue
 - May unfairly get in front of blocked processes