



University of
Pittsburgh

Introduction to Operating Systems

CS 1550



Spring 2023

Sherif Khattab

ksm73@pitt.edu

(Some slides are from **Silberschatz, Galvin and Gagne ©2013**)

Announcements

- Upcoming deadlines
 - Homework 6 is due **this Friday**
 - Quiz 1 and Lab 2 due on Tuesday 2/28 at 11:59 pm
 - Project 2 is due Friday 3/17 at 11:59 pm
- Midterm exam on Thursday 3/2
 - In-person, on paper, closed book
 - Study guide, old exam, and practice Midterm on Canvas
- Lost points because autograder or simple mistake?
 - please reach out to Grader TA over Piazza
- Navigating the Panopto Videos
 - Video contents
 - Search in captions

Previous lecture ...

- Sleepy Barbers solution using Condition Variables

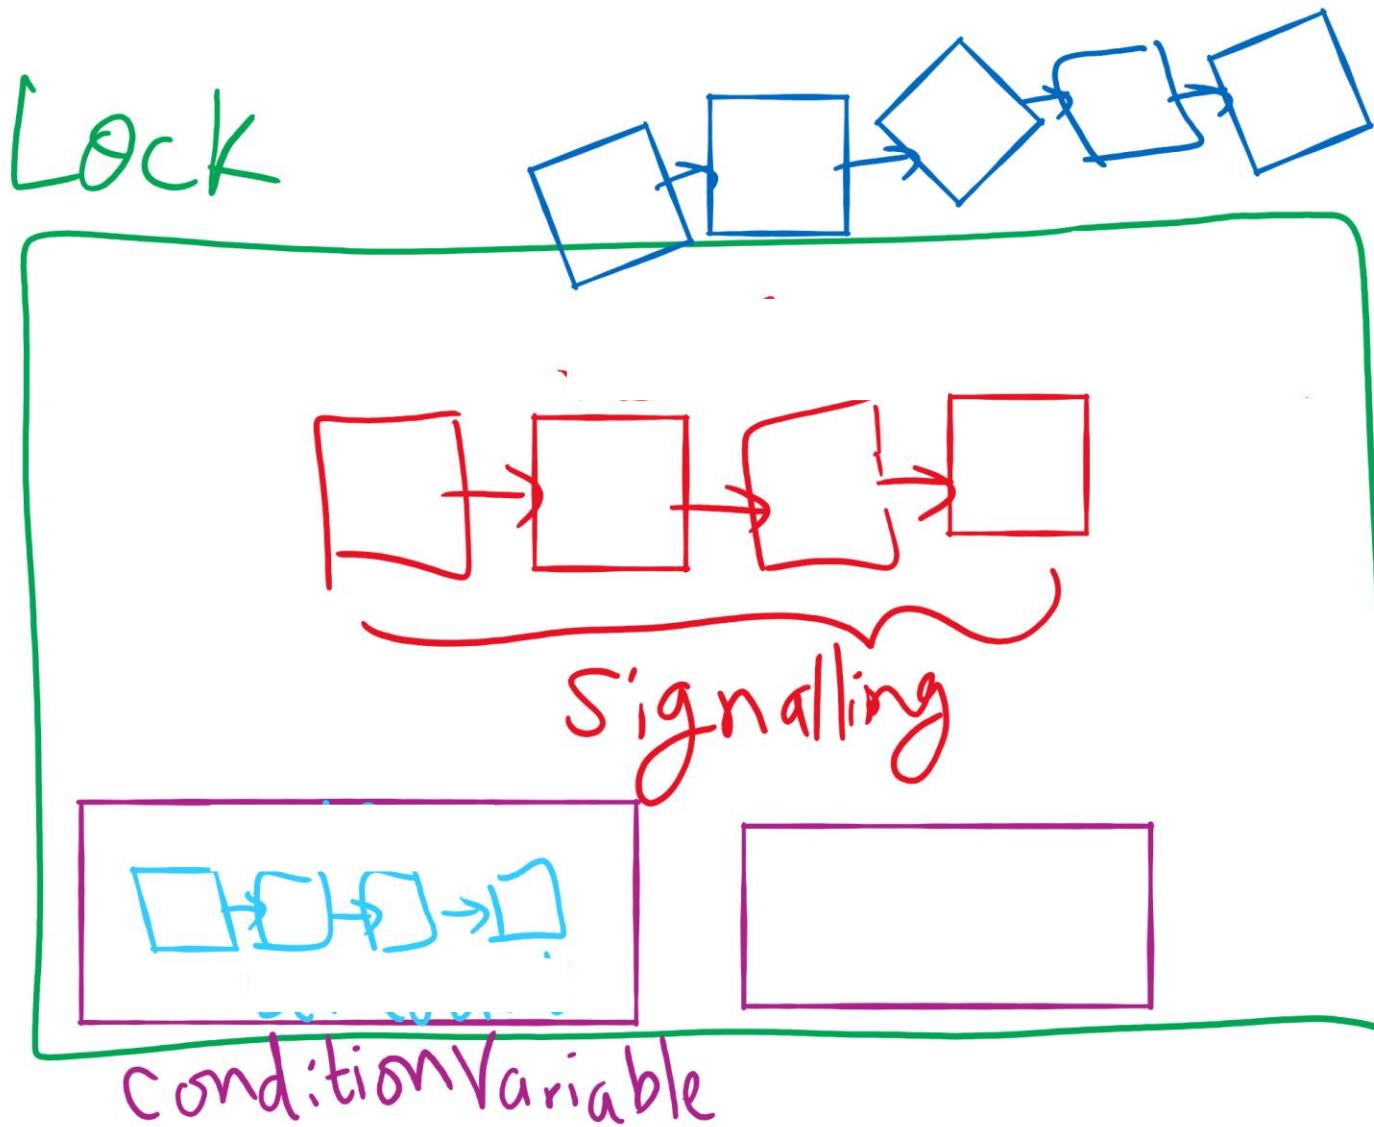
Today ...

- How to implement condition variables
- Reflections on using semaphores and condition variables
- CPU Scheduling

User-level implementation of Condition Variables

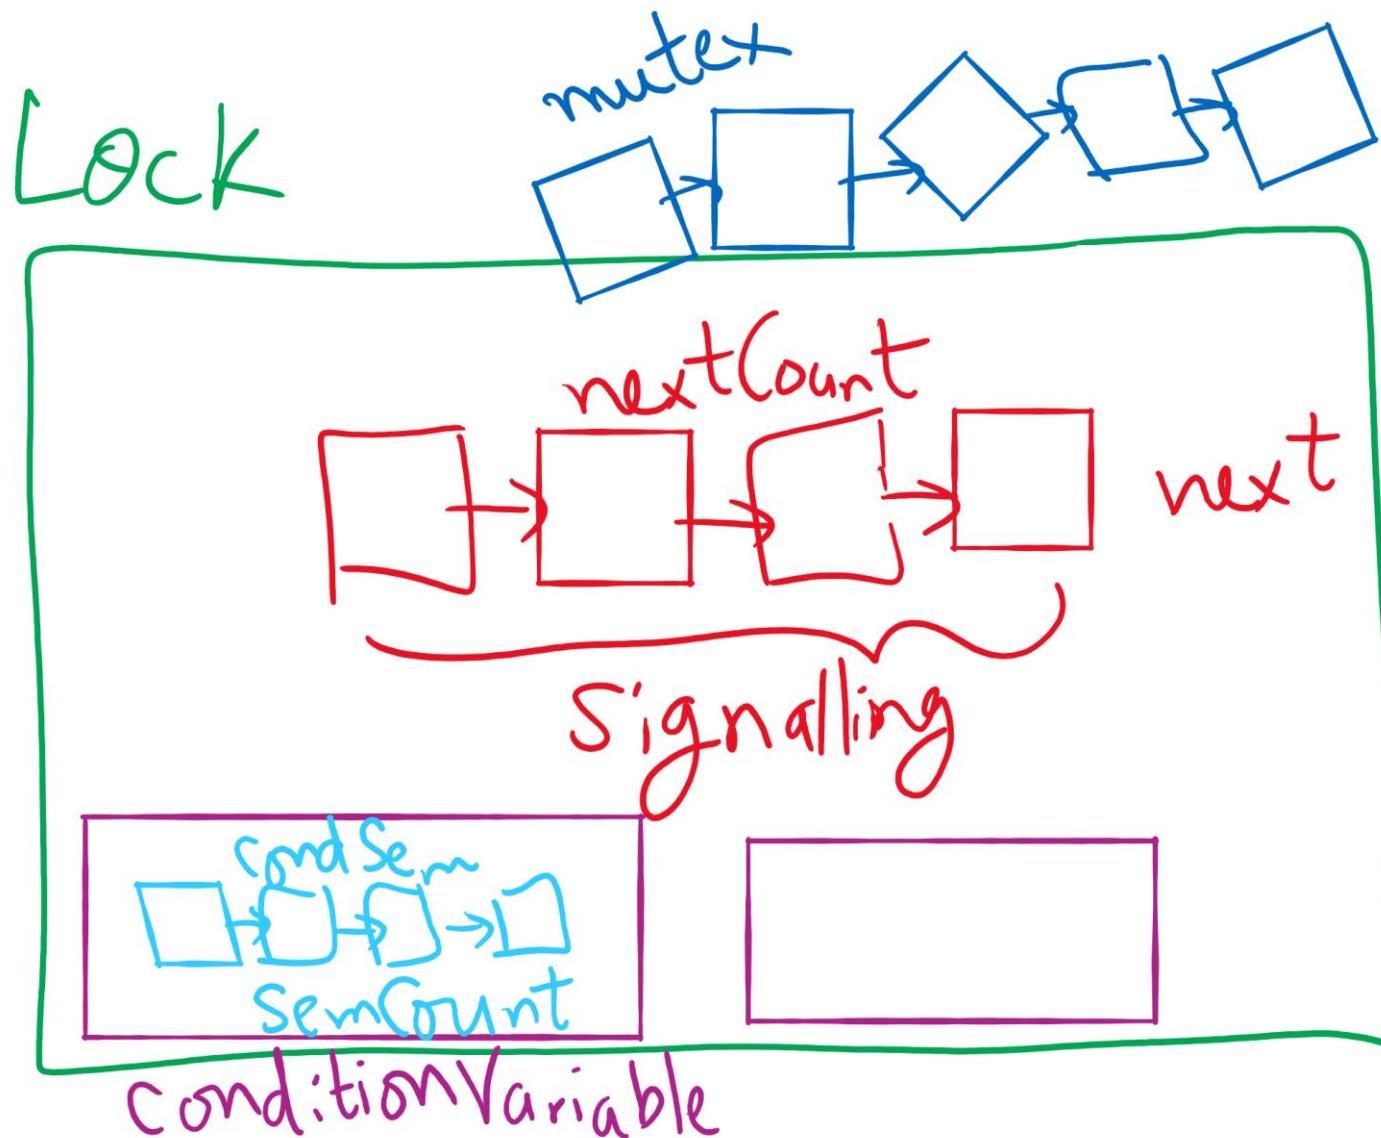
- Why?
 - Some operating systems don't have condition variables
 - Another exercise on solving synchronization problems with semaphores
- What are the waiting situations?
 - waiting to acquire the lock
 - waiting on the condition variable
 - waiting after signaling on a condition variable
 - This is Hoare semantics
 - signaling process waits
 - Compare to Mesa semantics
 - signaled process waits

Lock and Condition Variable Implementation



Lock and Condition Variable Implementation

- Let's have a semaphore for each waiting situation



User-level implementation of Condition Variables

A Lock with two waiting queues

```
struct Lock {  
    Semaphore mutex(1);  
    Semaphore next(0);  
    int nextCount = 0;  
}
```

```
Acquire(){  
    mutex.down();  
}
```

```
Release(){  
    if(nextCount > 0){  
        next.up();  
        nextCount--;  
    } else mutex.up();  
}
```

Condition Variable

```
struct ConditionVariable {  
    Semaphore condSem(0);  
    int semCount = 0;  
    Lock *lk;  
}
```

```
Wait(){
```

```
    if(lk->nextCount > 0)
```

```
        lk->next.up();
```

```
        lk->nextCount--;
```

```
    else lk->mutex.up();
```

```
    semCount++;
```

```
    condSem.down();
```

```
    semCount--;
```

```
}
```

```
Signal(){
```

```
    if(semCount > 0){
```

```
        condSem.up()
```

```
        lk->nextCount++
```

```
        lk->next.down();
```

```
        lk->nextCount—
```

```
}
```

```
}
```

Let's trace our solution!

- Note: Monitor is another name for Lock

P
≡
enter monitor
≡
≡ while (cond, tim)
CV.Wait()
≡
exit monitor



Q
≡
≡
enter monitor
≡
≡ change variables to make condition false
CV.Signal()
≡
exit monitor



Reflections on semaphore usage

- Semaphores can be used as
 - Resource counters
 - Waiting spaces
 - For mutual exclusion

Reflections on Condition Variables

- Define a class and put all shared variables inside the class
- Include a mutex and a condition variable in the class
- For each public method of the class
 - Start by locking the mutex lock
 - If need to wait, use a while loop and wait on the condition variable
 - Before broadcasting on the condition variable, make sure to change the waiting condition

Final Remarks on Process Synchronization

- Many other synchronization mechanisms
 - Message passing
 - Barriers
 - Futex
 - Re-entrant locks
 - AtomicInteger, AtomicX

Problem of the Day: CPU Scheduling

How does the ***short-term scheduler*** select the next process to run?

CPU Scheduling

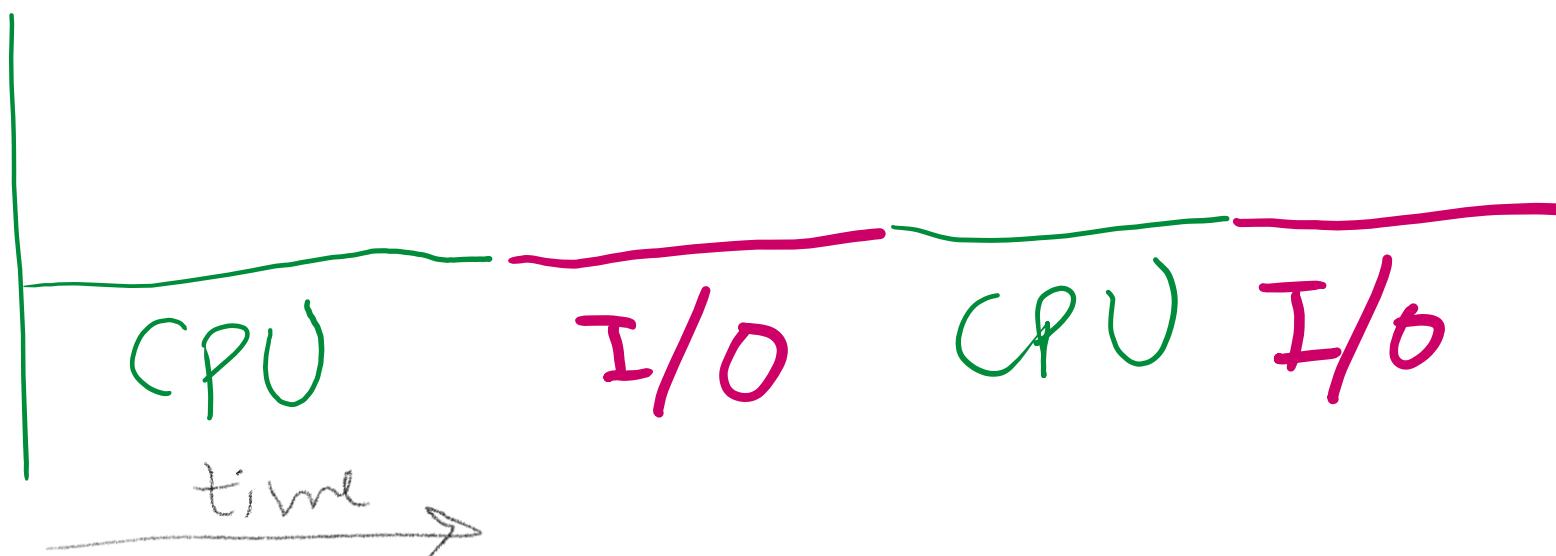
- Scheduling the processor among all ready processes
- User-oriented criteria
 - **Response Time**: Elapsed time between the submission of a request and the receipt of a response
 - **Turnaround Time**: Elapsed time between the submission of a process to its completion
- System-oriented criteria
 - Processor utilization
 - Throughput: number of process completed per unit time
 - Fairness

Short-Term Scheduler Dispatcher

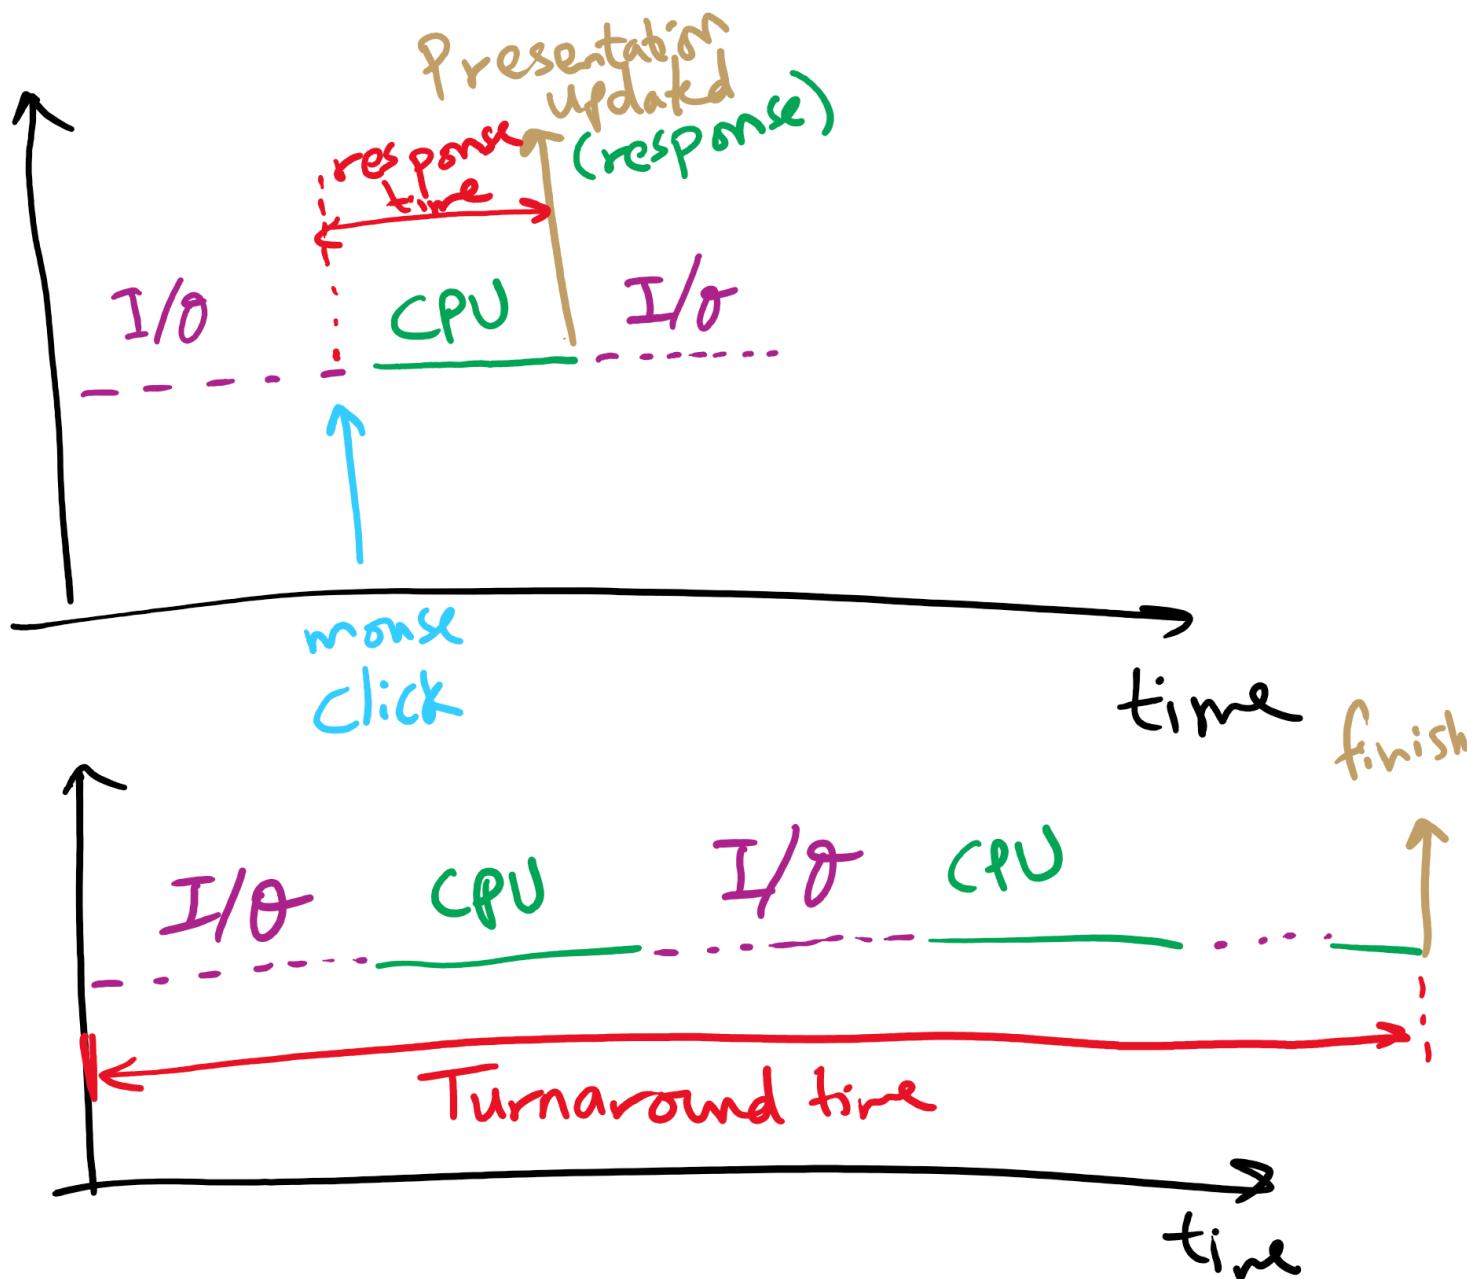
- The dispatcher is the module that gives control of the CPU to the process selected by the short-term scheduler
- The functions of the dispatcher include:
 - Switching context
 - Switching to user mode
 - Jumping to the location in the user program to restart execution
- The dispatch latency must be minimal

The CPU-I/O Cycle

- Processes require alternate use of processor and I/O in a repetitive fashion
- Each cycle consist of a CPU burst followed by an I/O burst
 - A process terminates on a CPU burst
- CPU-bound processes have longer CPU bursts than I/O-bound processes



Response time vs. Turnaround time



Scheduling Algorithms

- First-Come, First-Served Scheduling
- Shortest-Job-First Scheduling
 - Also referred to as Shortest Process Next
- Priority Scheduling
- Round-Robin Scheduling
- Multilevel Queue Scheduling
- Multilevel Feedback Queue Scheduling

Characterization of Scheduling Policies

- The **selection function** determines which ready process is selected next for execution
- The **decision mode** specifies the instants in time the selection function is exercised
 - Nonpreemptive
 - Once a process is in the running state, it will continue until it terminates or blocks for an I/O
 - Preemptive
 - Currently running process may be interrupted and moved to the Ready state by the OS
 - Prevents one process from monopolizing the processor

Process Mix Example

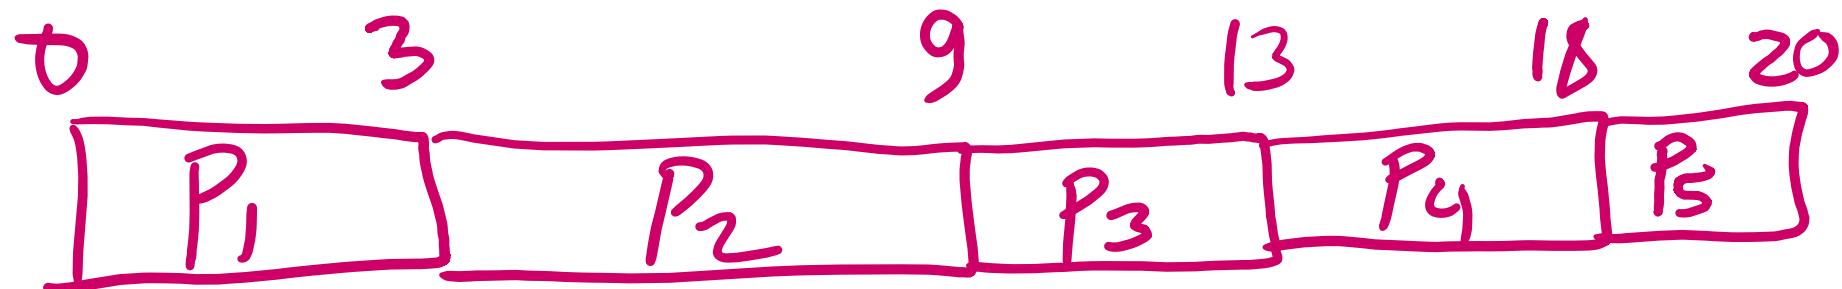
| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |

Service time = total processor time needed in one (CPU-I/O) cycle
Jobs with long service time are CPU-bound jobs and are referred to as “long jobs”

First Come First Served (FCFS)

- Selection function: the process that has been waiting the longest in the ready queue (hence, FCFS)
- Decision mode: non-preemptive
 - a process runs until it blocks for an I/O

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



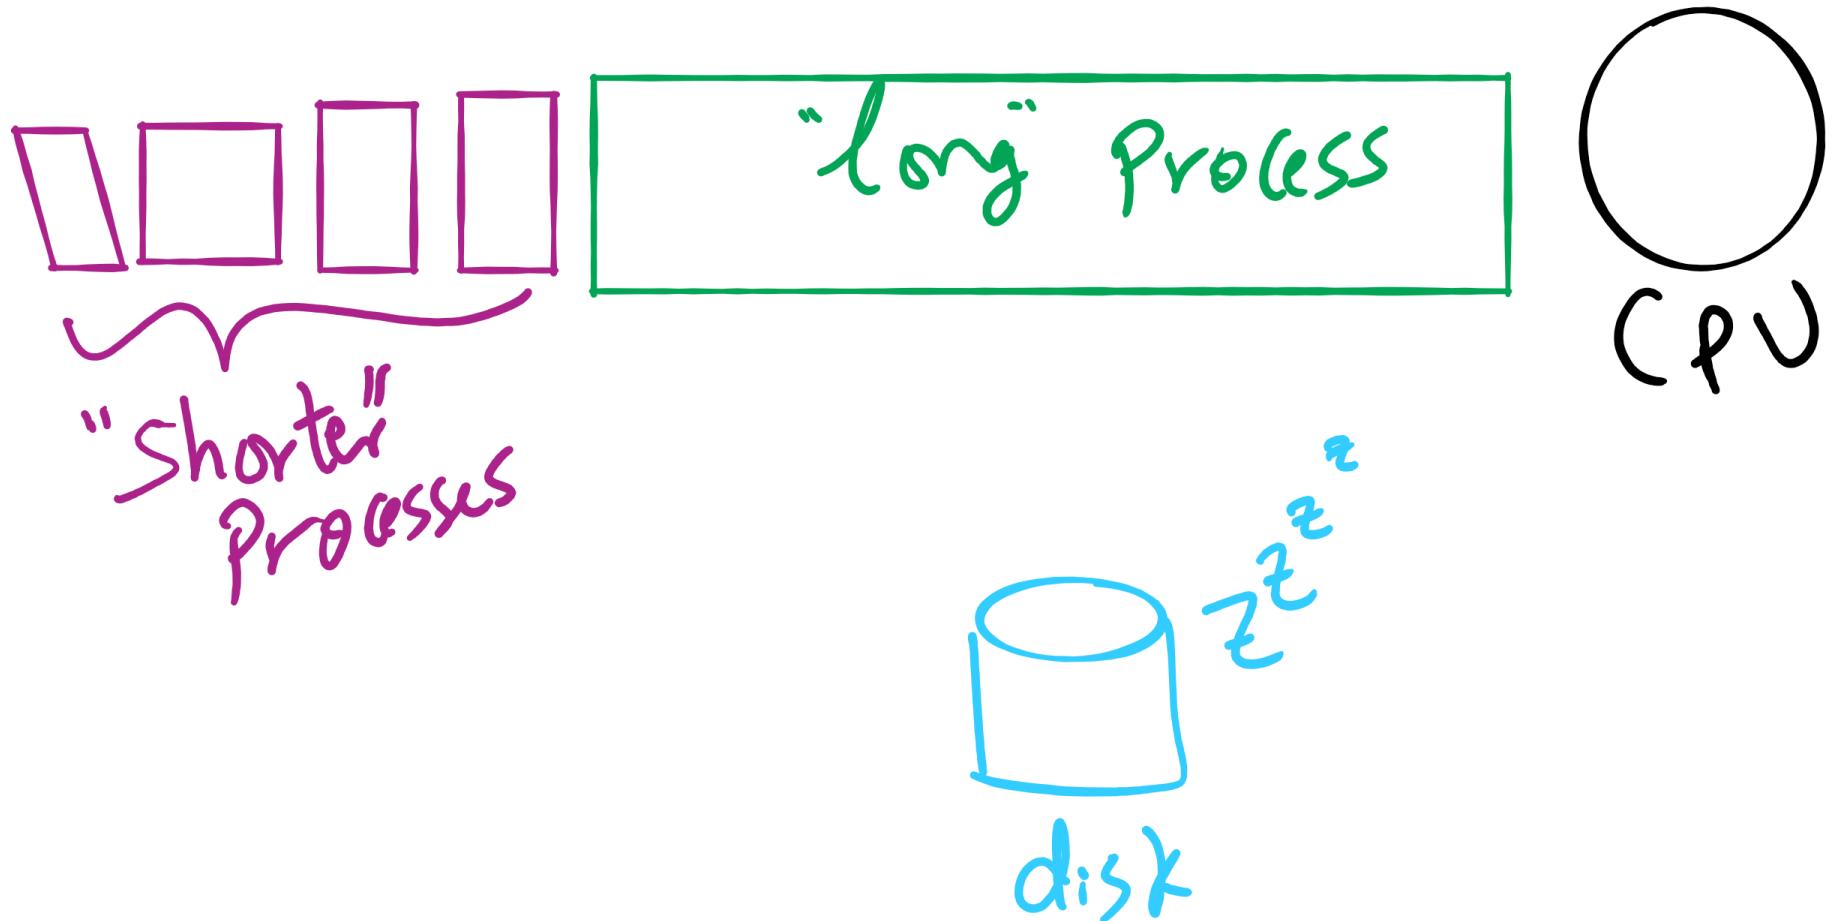
Average Response Time

$$\text{Average Response Time} = \frac{(3-0) + (9-2) + (13-4) + (18-6) + (20-8)}{5}$$

FCFS drawbacks

- Favours CPU-bound processes
 - CPU-bound processes monopolize the processor
 - I/O-bound processes have to wait until completion of CPU-bound process
 - I/O-bound processes may have to wait even after their I/Os are completed (poor device utilization)
 - Convoy effect
 - Better I/O device utilization could be achieved if I/O bound processes had higher priority

Convoy Effect



Shortest Job First (Shortest Process Next)

- Selection function: the process with the shortest expected CPU burst time
 - I/O-bound processes will be selected first
- Decision mode: non-preemptive
- The required processing time, i.e., the CPU burst time, must be estimated for each process

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



SJF / SPN Critique

- Possibility of starvation for longer processes
- Lack of preemption is not suitable in a time sharing environment
- SJF/SPN implicitly incorporates priorities
 - Shortest jobs are given preference
 - CPU bound processes have lower priority, but a process doing no I/O could still monopolize the CPU if it is the first to enter the system

Priorities

- Implemented by having multiple ready queues to represent each level of priority
- Scheduler selects the process of a higher priority over one of lower priority
- Lower-priority may suffer starvation
- To alleviate starvation allow dynamic priorities
 - The priority of a process changes based on its age or execution history

Round-Robin

- Selection function: same as FCFS
- Decision mode: pre-emptive

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |

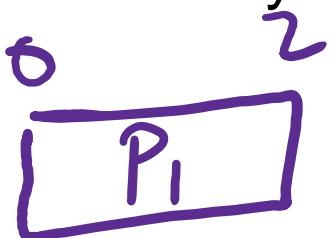
- a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired
- a clock interrupt occurs and the running process is put on the ready queue

Round-Robin

- Selection function: same as FCFS
- Decision mode: pre-emptive

- a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- a clock interrupt occurs and the running process is put on the ready queue



| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |

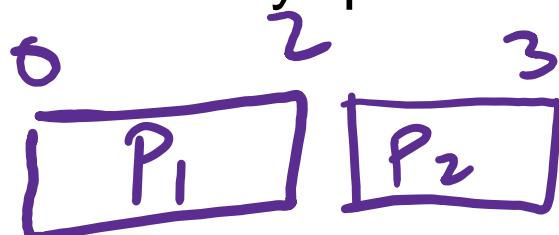
Round-Robin

- Selection function: same as FCFS

- Decision mode: pre-emptive

- a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- a clock interrupt occurs and the running process is put on the ready queue



| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |

Round-Robin

- Selection function: same as FCFS

- Decision mode: pre-emptive

- a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- a clock interrupt occurs and the running process is put on the ready queue

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Round-Robin

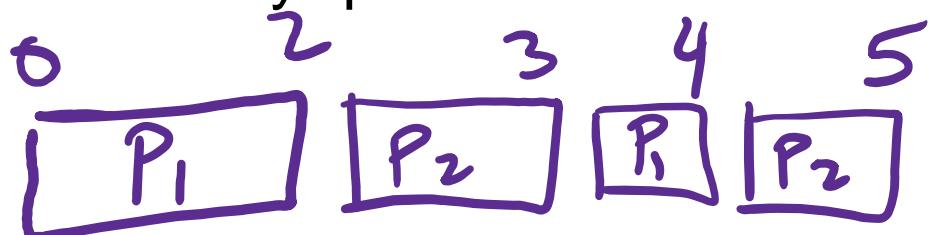
- Selection function: same as FCFS

- Decision mode: pre-emptive

- a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- a clock interrupt occurs and the running process is put on the ready queue

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Round-Robin

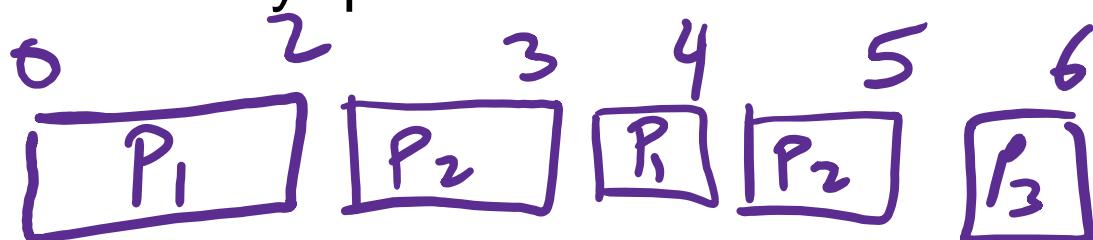
- Selection function: same as FCFS

- Decision mode: pre-emptive

- a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- a clock interrupt occurs and the running process is put on the ready queue

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Round-Robin

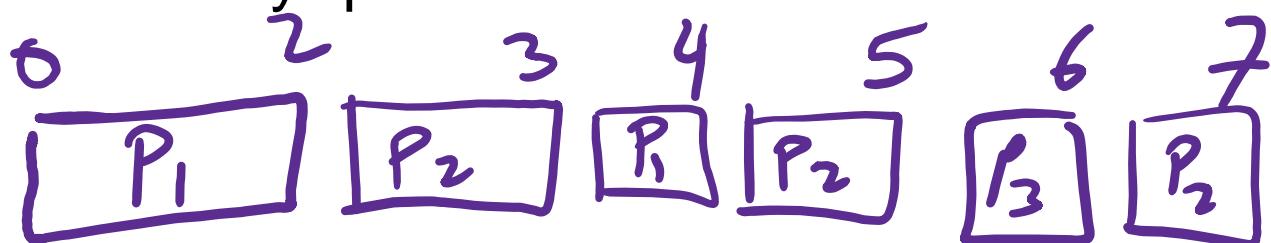
- Selection function: same as FCFS

- Decision mode: pre-emptive

- a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- a clock interrupt occurs and the running process is put on the ready queue

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Round-Robin

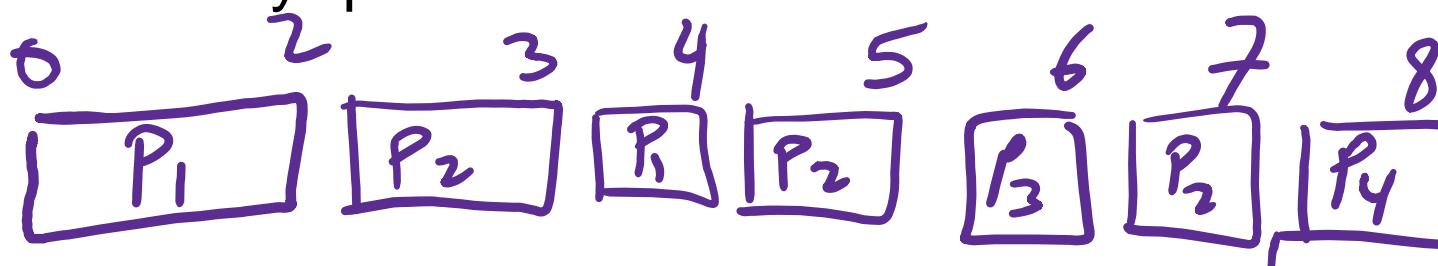
- Selection function: same as FCFS

- Decision mode: pre-emptive

- a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- a clock interrupt occurs and the running process is put on the ready queue

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Round-Robin

- Selection function: same as FCFS

- Decision mode: pre-emptive

- a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- a clock interrupt occurs and the running process is put on the ready queue

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Round-Robin

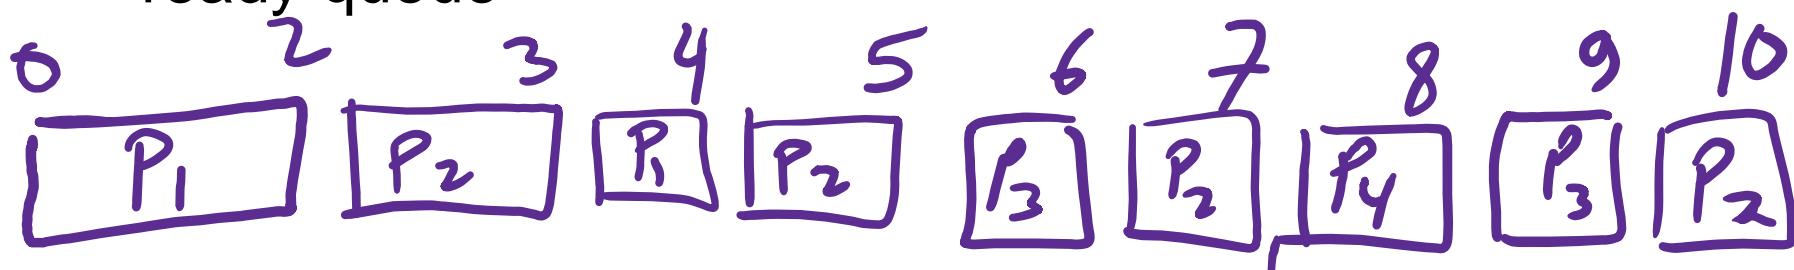
- Selection function: same as FCFS

- Decision mode: pre-emptive

- a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- a clock interrupt occurs and the running process is put on the ready queue

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Round-Robin

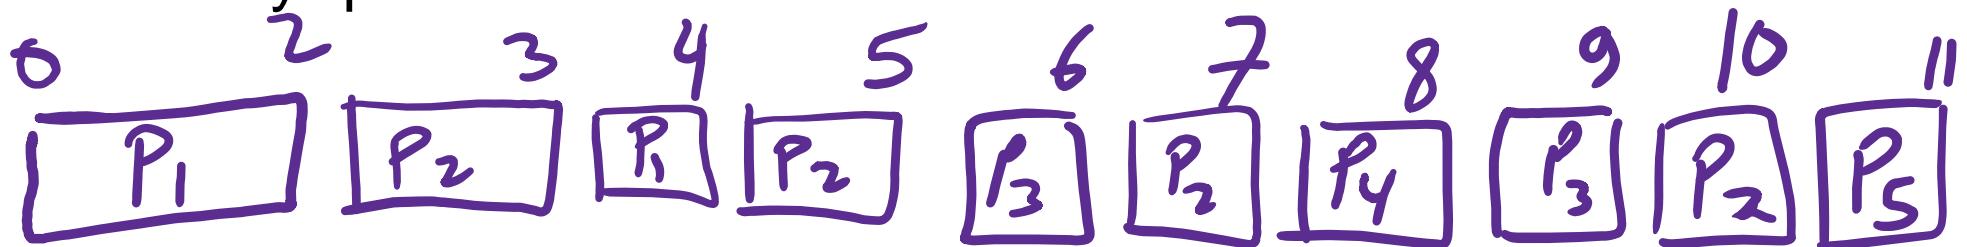
- Selection function: same as FCFS

- Decision mode: pre-emptive

- a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- a clock interrupt occurs and the running process is put on the ready queue

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Round-Robin

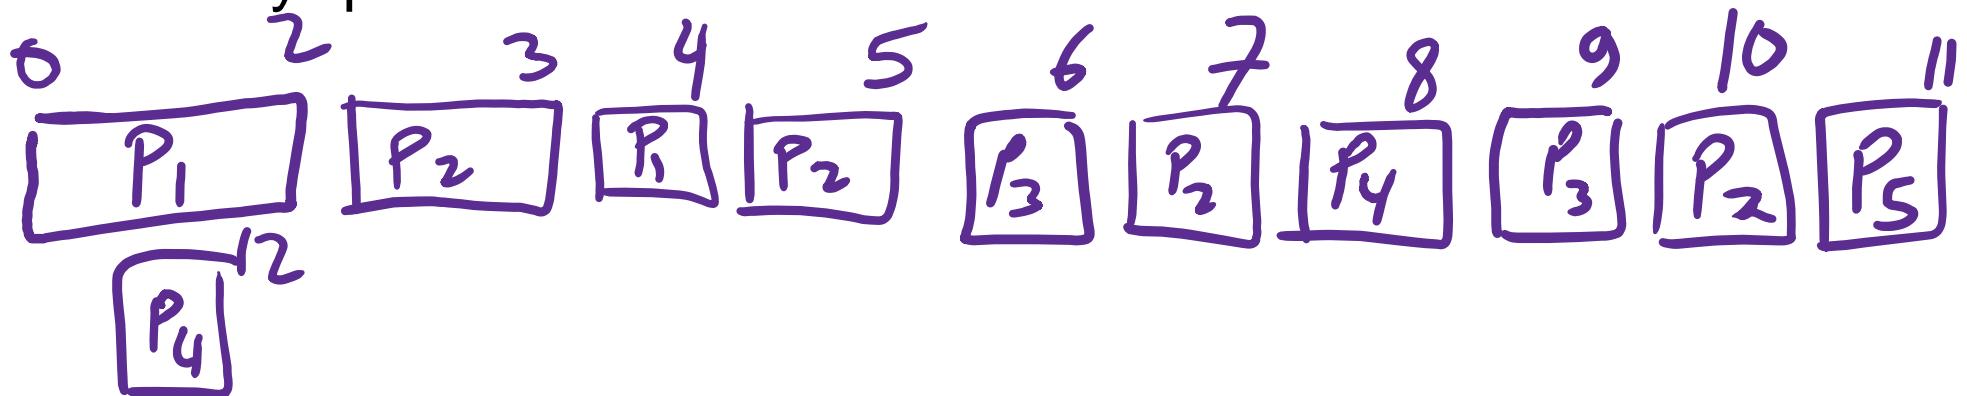
- Selection function: same as FCFS

- Decision mode: pre-emptive

- a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- a clock interrupt occurs and the running process is put on the ready queue

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Round-Robin

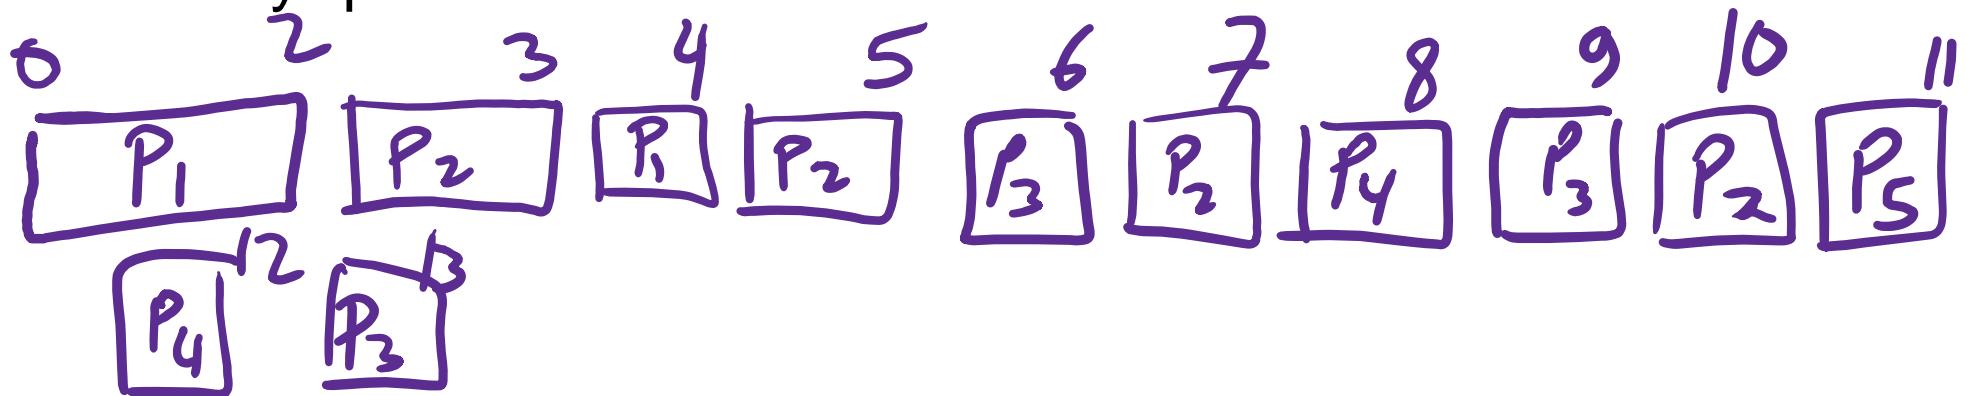
- Selection function: same as FCFS

- Decision mode: pre-emptive

- a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- a clock interrupt occurs and the running process is put on the ready queue

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Round-Robin

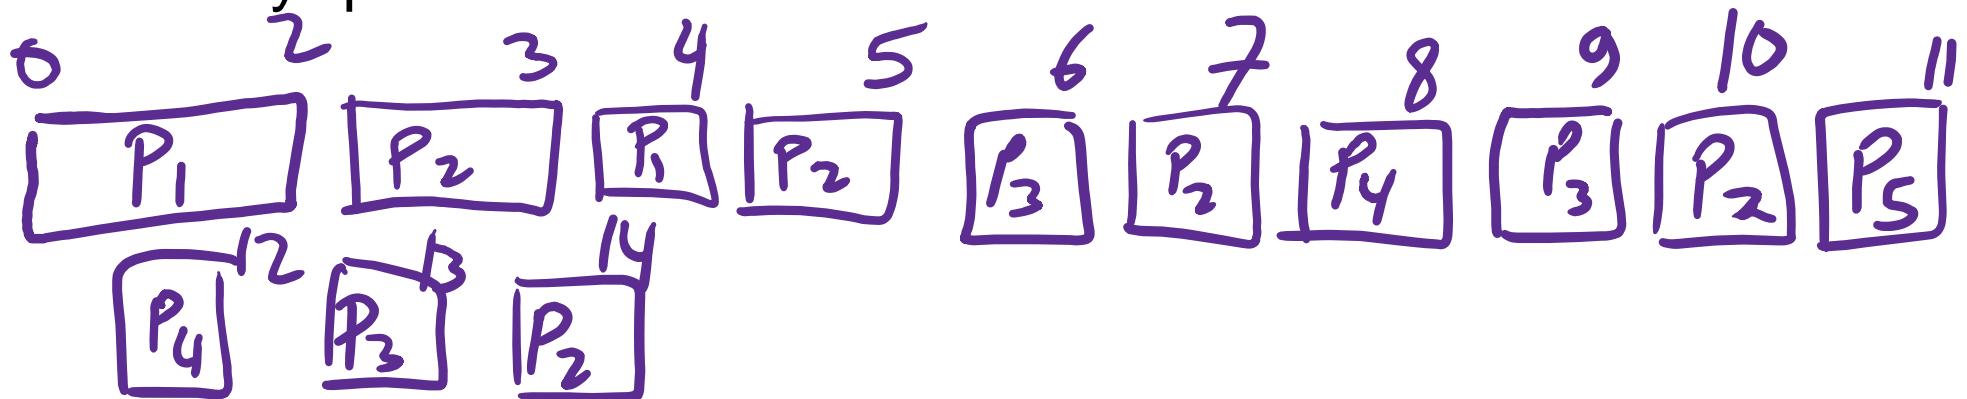
- Selection function: same as FCFS

- Decision mode: pre-emptive

- a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- a clock interrupt occurs and the running process is put on the ready queue

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Round-Robin

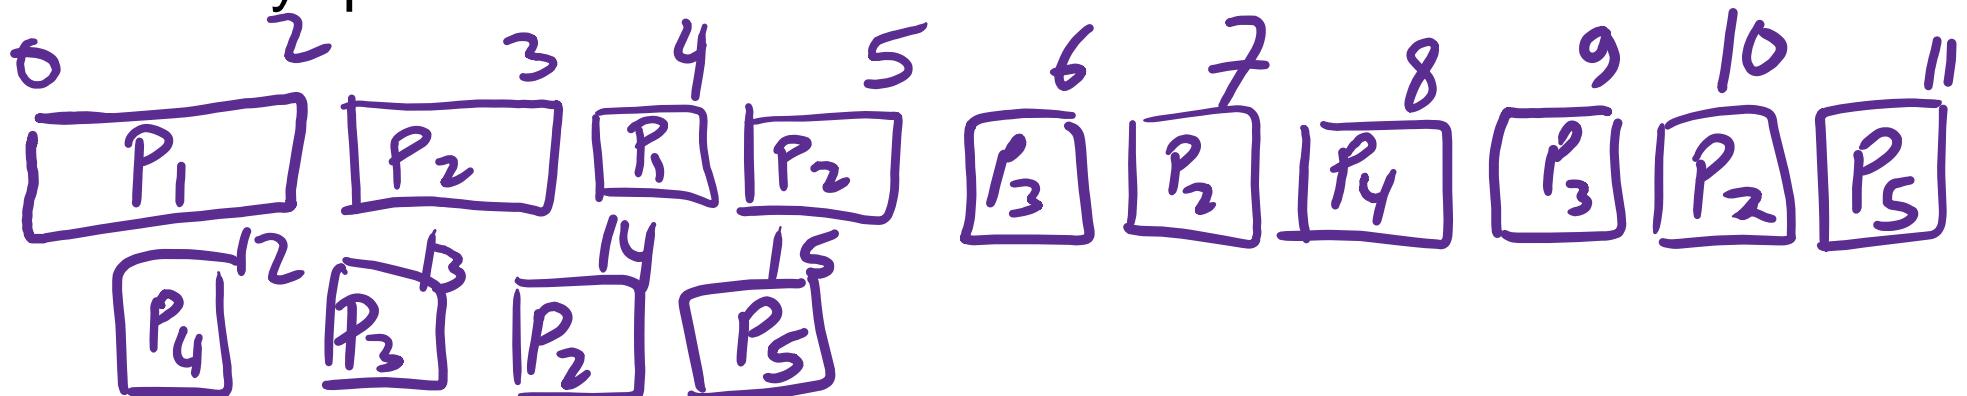
- Selection function: same as FCFS

- Decision mode: pre-emptive

- a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- a clock interrupt occurs and the running process is put on the ready queue

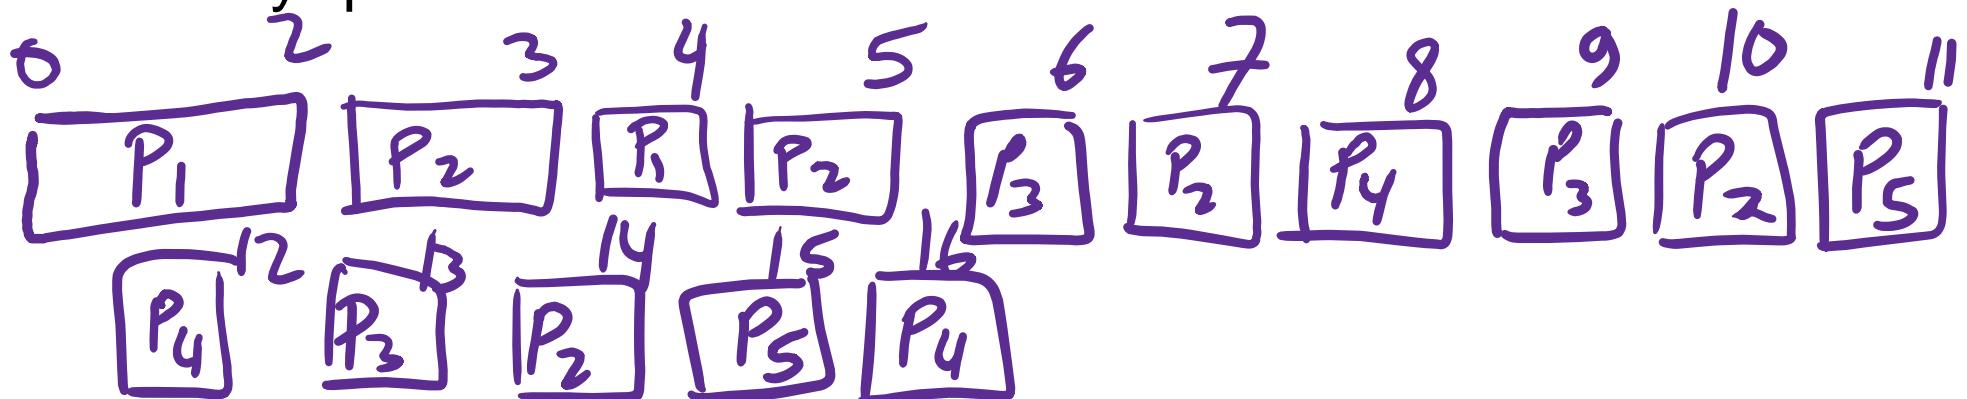
| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Round-Robin

- Selection function: same as FCFS
- Decision mode: pre-emptive
 - a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired
 - a clock interrupt occurs and the running process is put on the ready queue

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Round-Robin

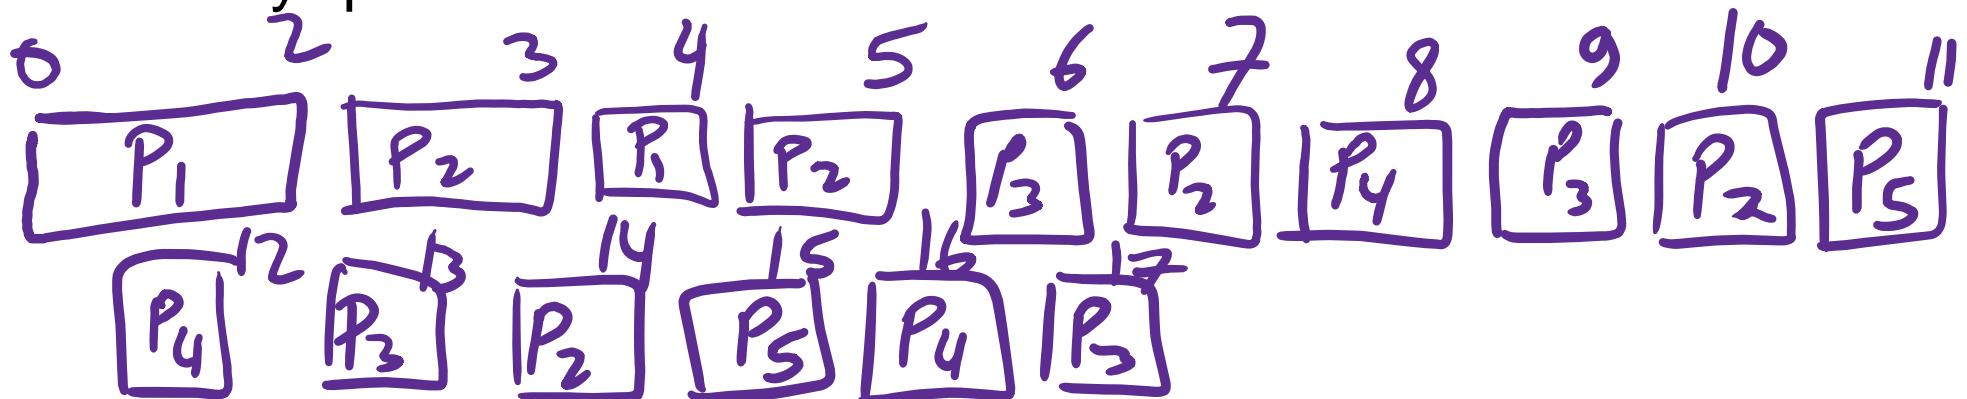
- Selection function: same as FCFS

- Decision mode: pre-emptive

- a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- a clock interrupt occurs and the running process is put on the ready queue

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Round-Robin

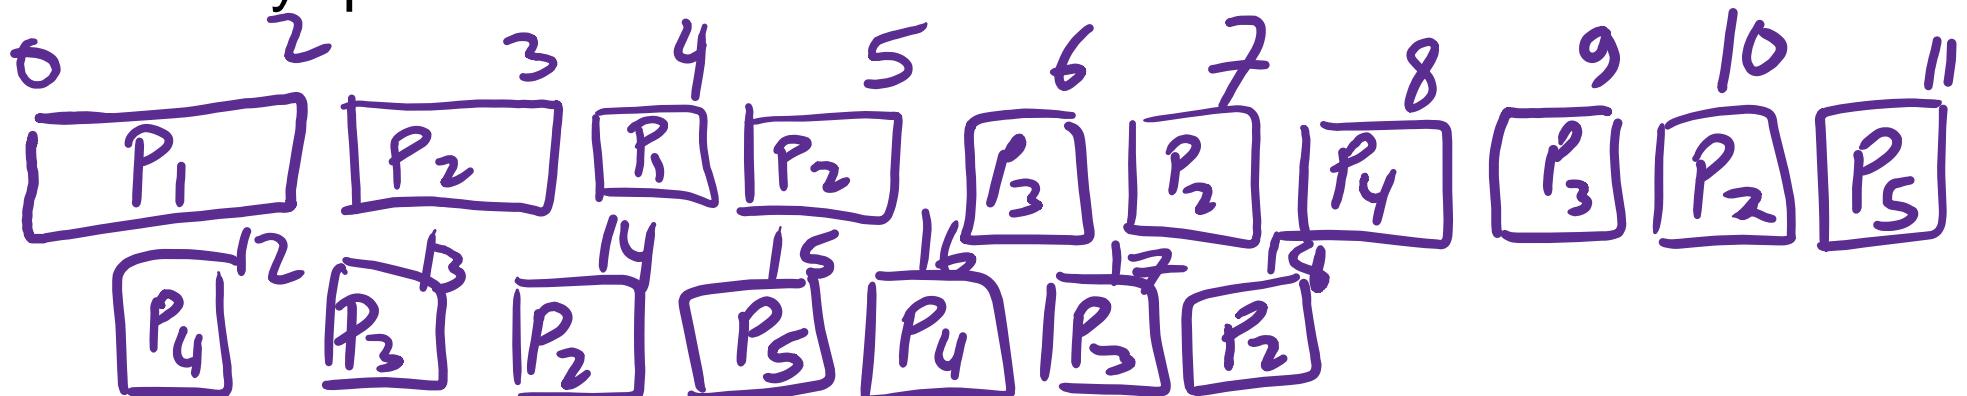
- Selection function: same as FCFS

- Decision mode: pre-emptive

- a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- a clock interrupt occurs and the running process is put on the ready queue

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Round-Robin

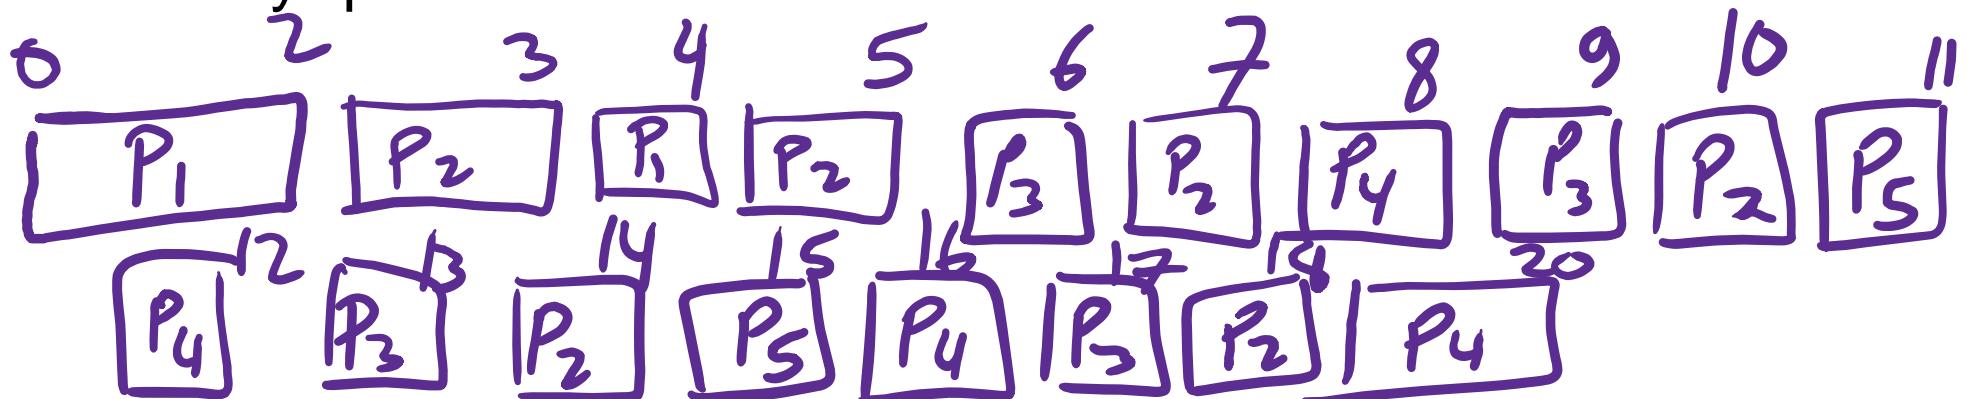
- Selection function: same as FCFS

- Decision mode: pre-emptive

- a process is allowed to run until the time slice period (quantum, typically from 10 to 100 ms) has expired

- a clock interrupt occurs and the running process is put on the ready queue

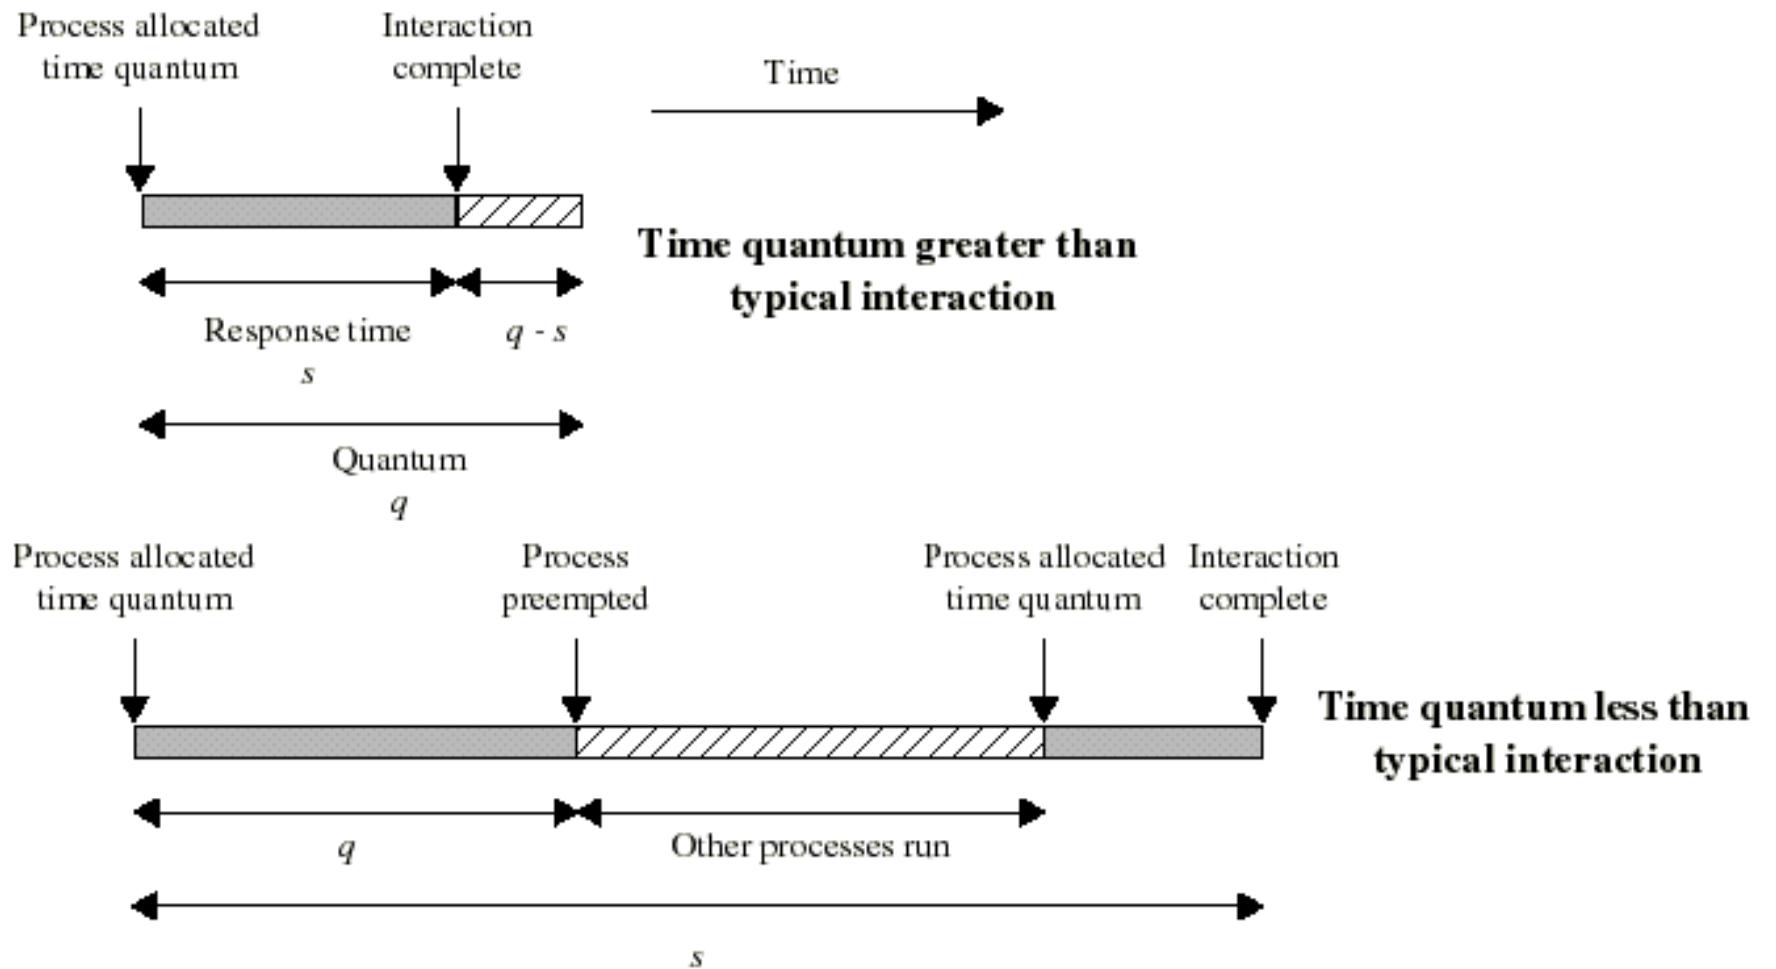
| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



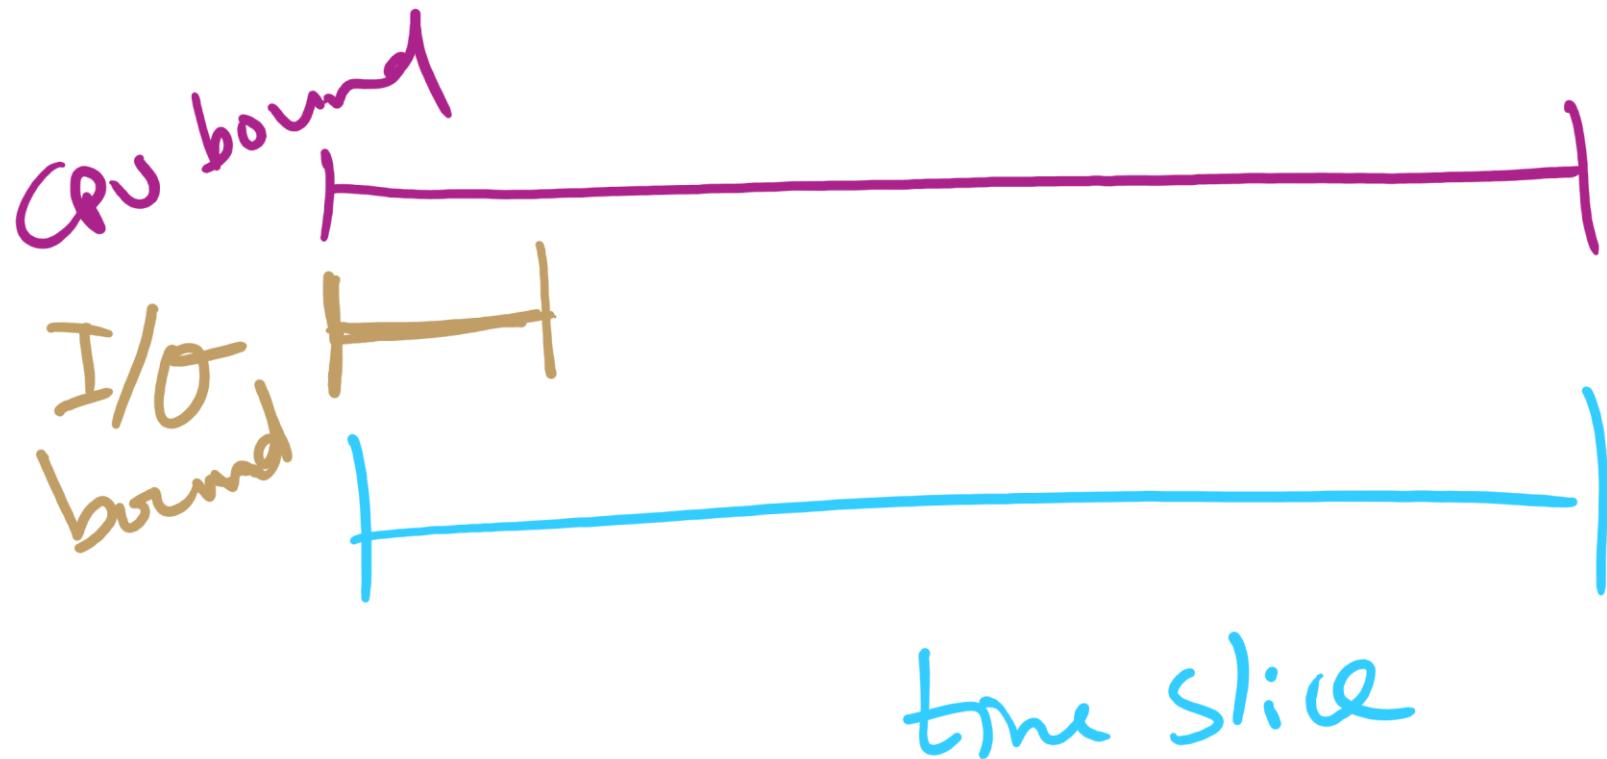
RR Time Quantum

- Quantum must be substantially larger than the time required to handle the clock interrupt and dispatching
- Quantum should be larger then the typical interaction
 - but not much larger, to avoid penalizing I/O bound processes

RR Time Quantum



Quantum Length



Round Robin: critique

- Still favors CPU-bound processes
 - An I/O bound process uses the CPU for a time less than the time quantum before it is blocked waiting for an I/O
 - A CPU-bound process runs for all its time slice and is put back into the ready queue
 - May unfairly get in front of blocked processes

Multilevel Feedback Scheduling

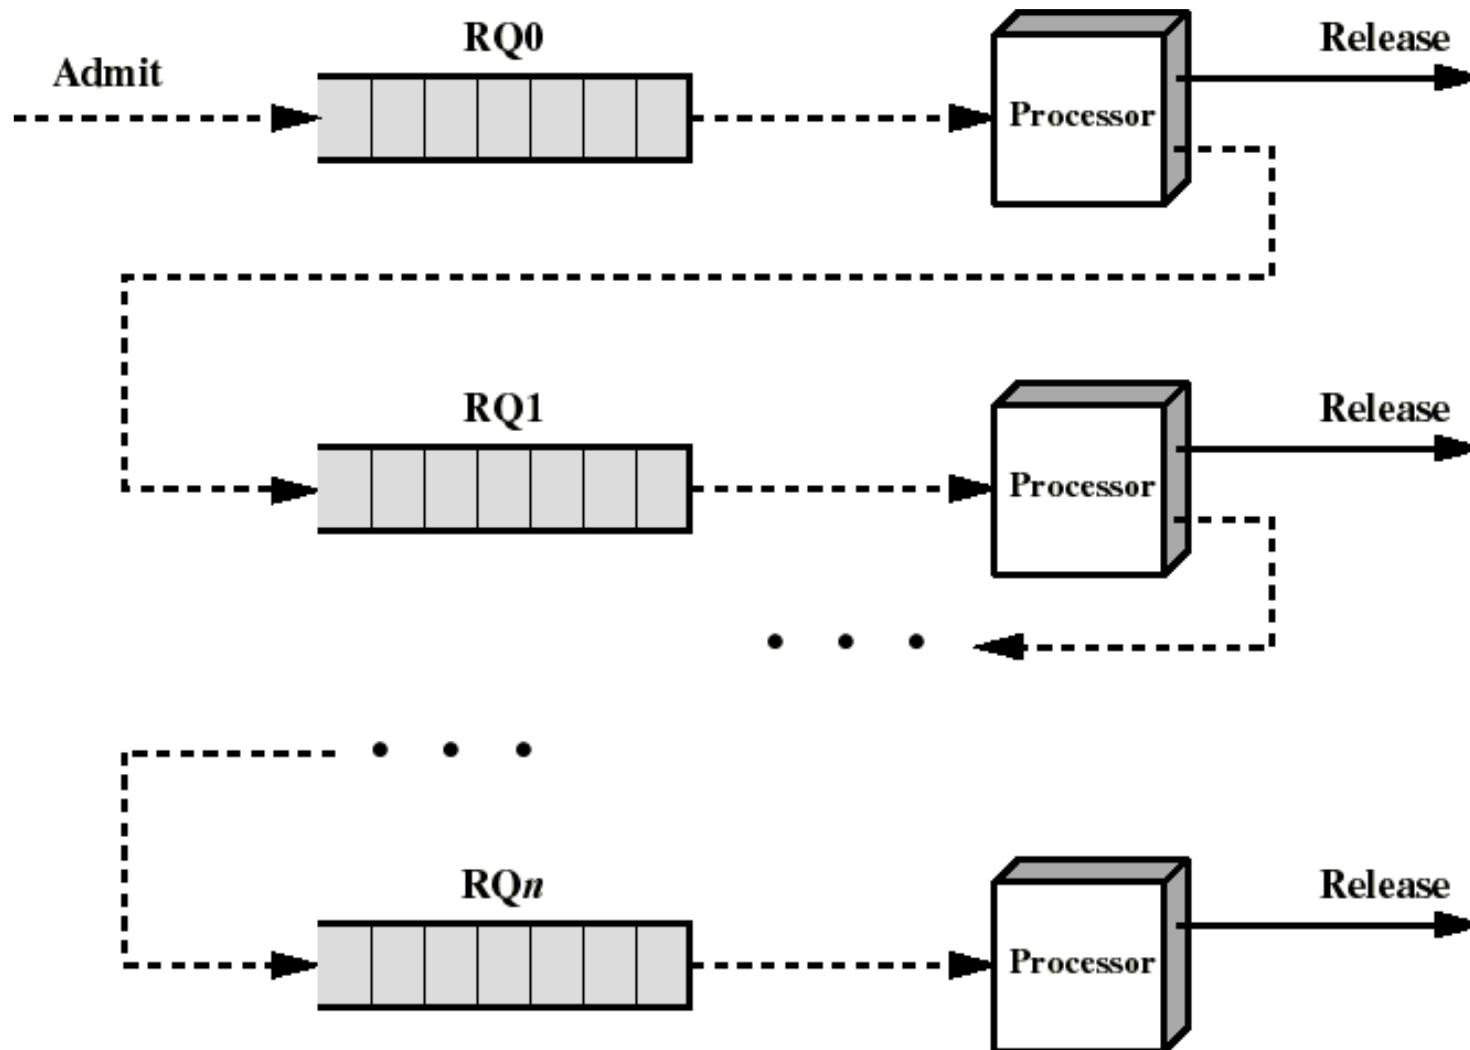
- Preemptive scheduling with dynamic priorities
- N ready to execute queues with decreasing priorities:
 - $P(RQ_0) > P(RQ_1) > \dots > P(RQ_{N-1})$
 - Dispatcher selects a process for execution from RQ_i only if RQ_{i-1} to RQ_0 are empty

Multilevel Feedback Scheduling

- New process are placed in RQ_0
- After the first quantum, they are moved to RQ_1 , and to RQ_2 after the second quantum, ... and to RQ_{N-1} after the Nth quantum
- I/O-bound processes remain in higher priority queues.
 - CPU-bound jobs drift downward.
 - Hence, long jobs may starve

Multiple Feedback Queues

Different RQs may have different quantum values



Time Quantum for feedback Scheduling

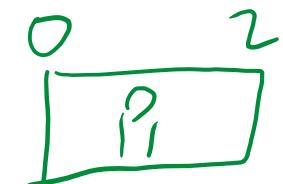
- With a fixed quantum time, the turn-around time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
 - Time quantum of $RQ_i = 2^{i-1}$
 - May still cause longer processes to suffer starvation.
 - Possible fix is to promote a process to higher queue after some time

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |

Time Quantum for feedback Scheduling

- With a fixed quantum time, the turn-around time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
 - Time quantum of $RQ_i = 2^{i-1}$
 - May still cause longer processes to suffer starvation.
 - Possible fix is to promote a process to higher queue after some time

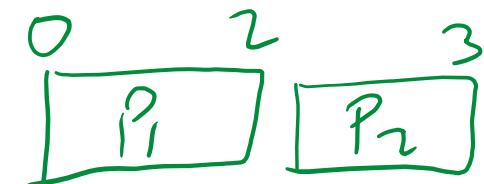
| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Time Quantum for feedback Scheduling

- With a fixed quantum time, the turn-around time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
 - Time quantum of $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
 - Possible fix is to promote a process to higher queue after some time

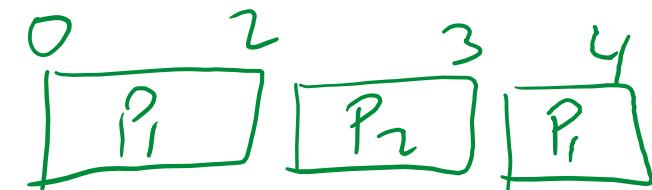
| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Time Quantum for feedback Scheduling

- With a fixed quantum time, the turn-around time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
 - Time quantum of $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
 - Possible fix is to promote a process to higher queue after some time

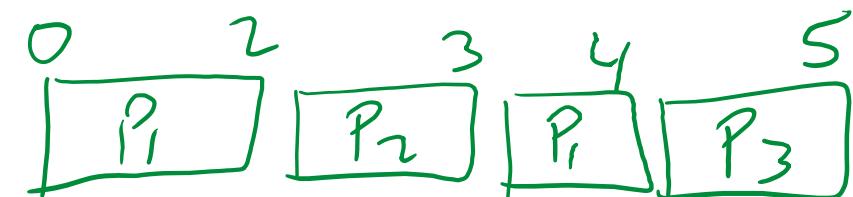
| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Time Quantum for feedback Scheduling

- With a fixed quantum time, the turn-around time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
 - Time quantum of $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
 - Possible fix is to promote a process to higher queue after some time

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Time Quantum for feedback Scheduling

- With a fixed quantum time, the turn-around time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
 - Time quantum of $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
 - Possible fix is to promote a process to higher queue after some time

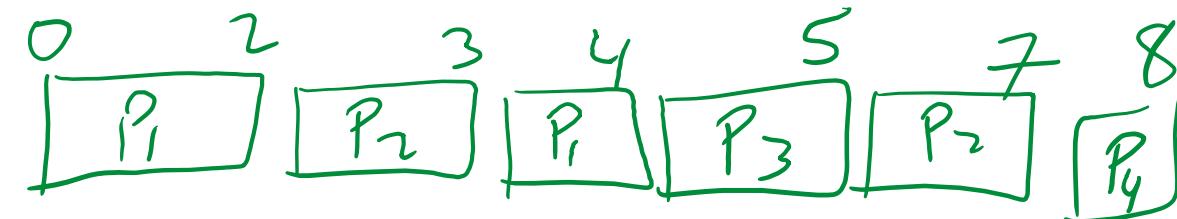
| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Time Quantum for feedback Scheduling

- With a fixed quantum time, the turn-around time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
 - Time quantum of $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
 - Possible fix is to promote a process to higher queue after some time

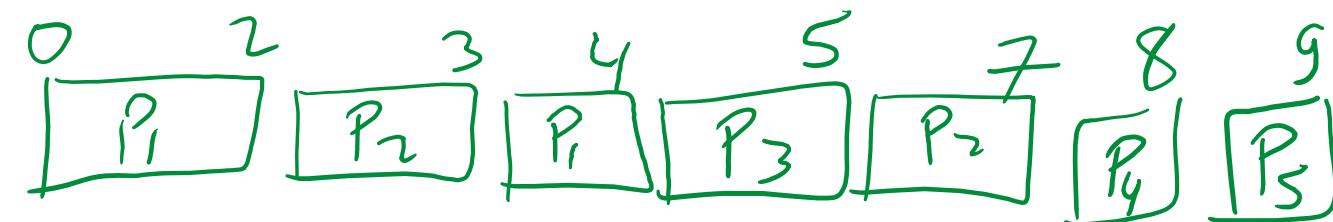
| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Time Quantum for feedback Scheduling

- With a fixed quantum time, the turn-around time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
 - Time quantum of $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
 - Possible fix is to promote a process to higher queue after some time

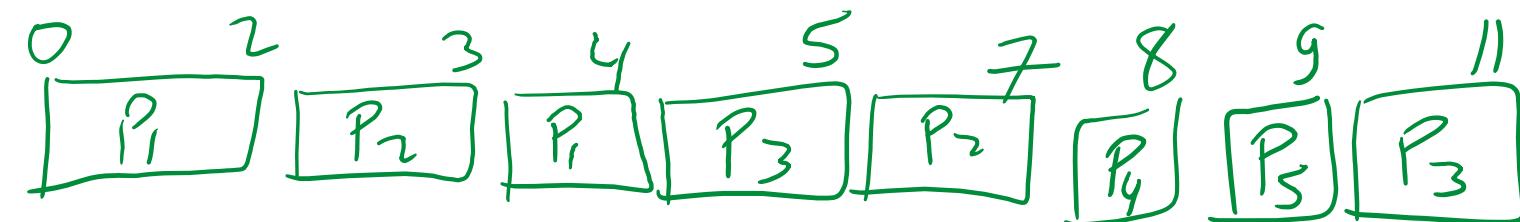
| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Time Quantum for feedback Scheduling

- With a fixed quantum time, the turn-around time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
 - Time quantum of $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
 - Possible fix is to promote a process to higher queue after some time

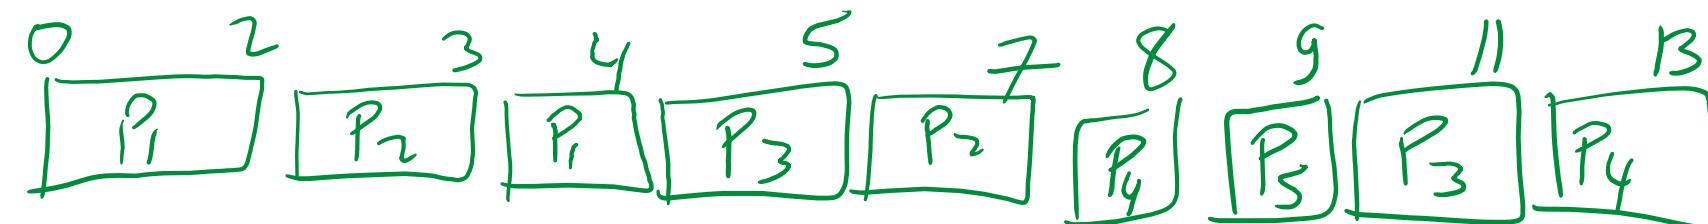
| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Time Quantum for feedback Scheduling

- With a fixed quantum time, the turn-around time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
 - Time quantum of $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
 - Possible fix is to promote a process to higher queue after some time

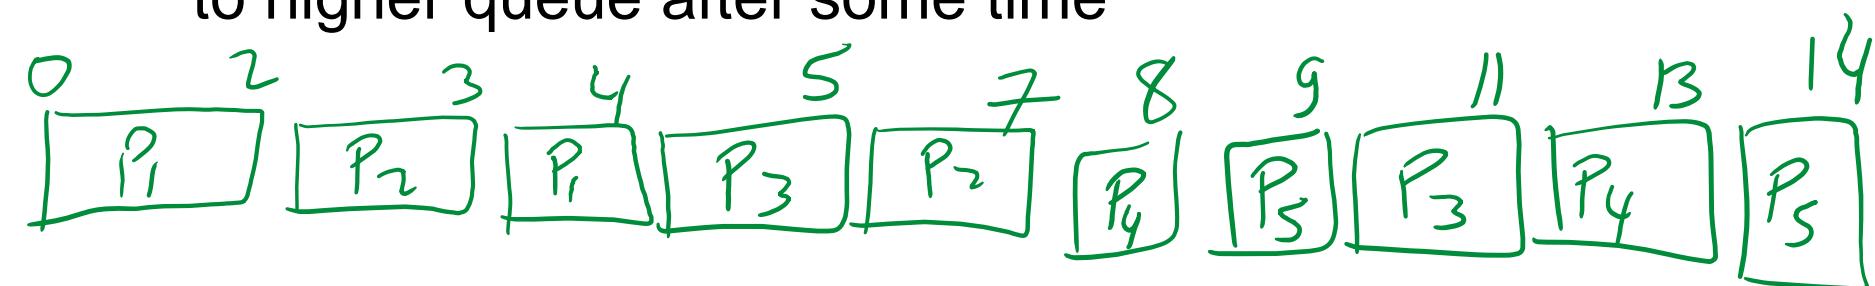
| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Time Quantum for feedback Scheduling

- With a fixed quantum time, the turn-around time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
 - Time quantum of $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
 - Possible fix is to promote a process to higher queue after some time

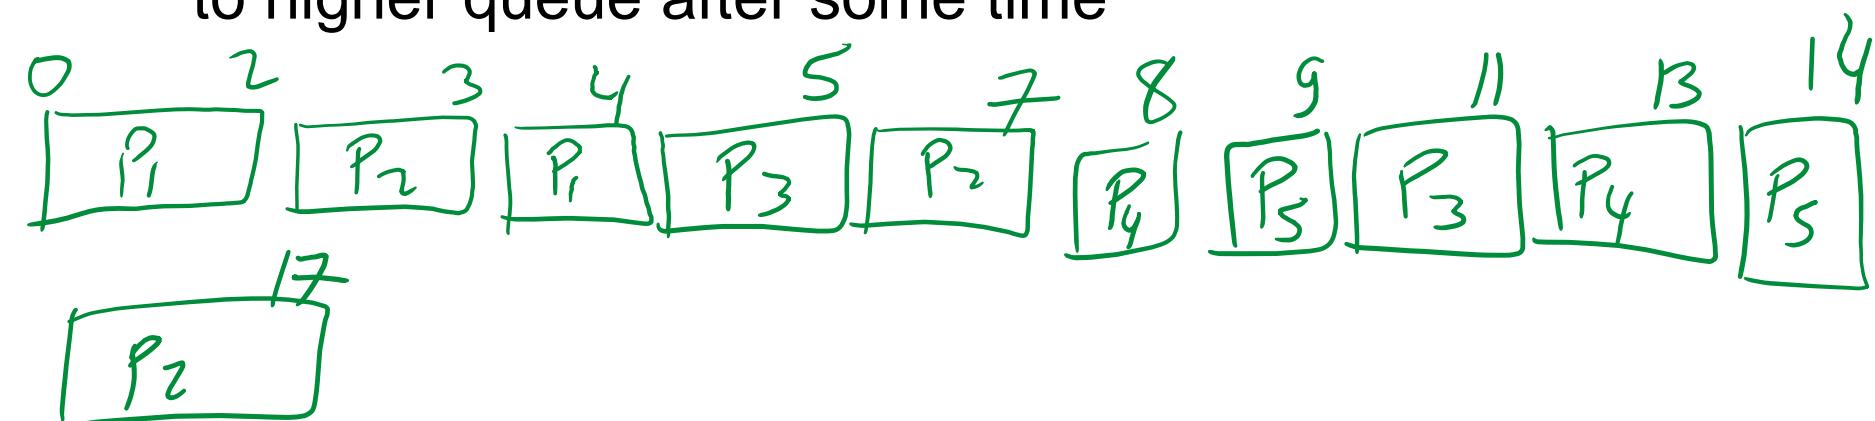
| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Time Quantum for feedback Scheduling

- With a fixed quantum time, the turn-around time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
 - Time quantum of $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
 - Possible fix is to promote a process to higher queue after some time

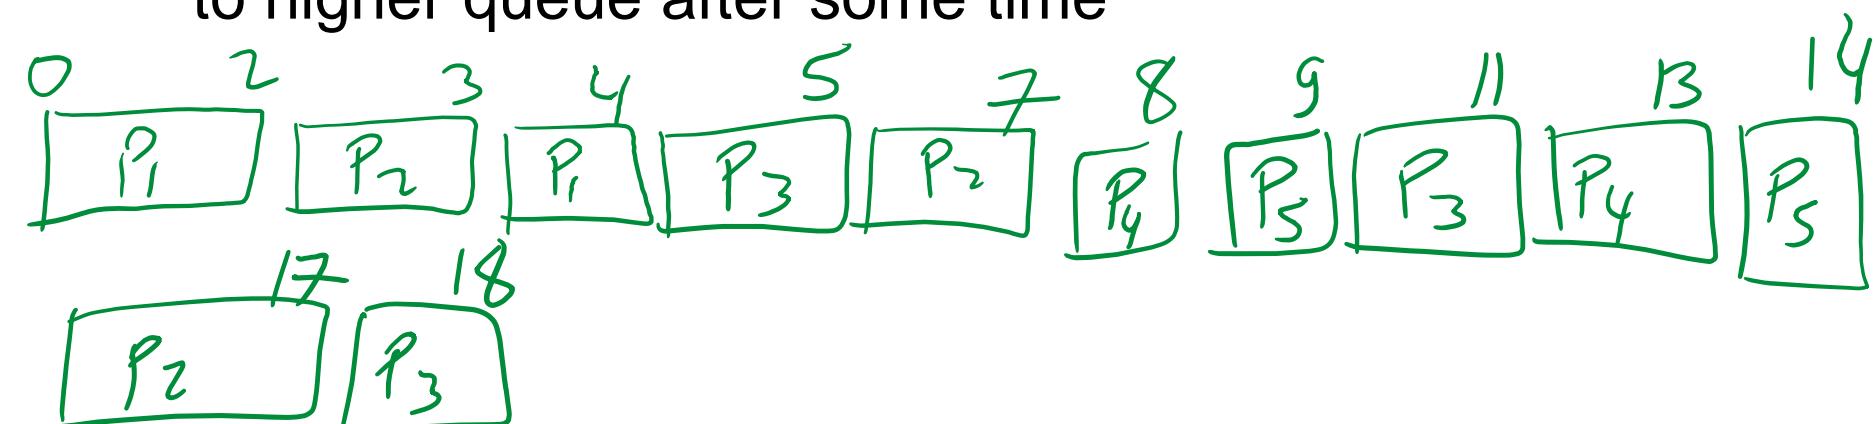
| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Time Quantum for feedback Scheduling

- With a fixed quantum time, the turn-around time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
 - Time quantum of $RQ_i = 2^{i-1}$
- May still cause longer processes to suffer starvation.
 - Possible fix is to promote a process to higher queue after some time

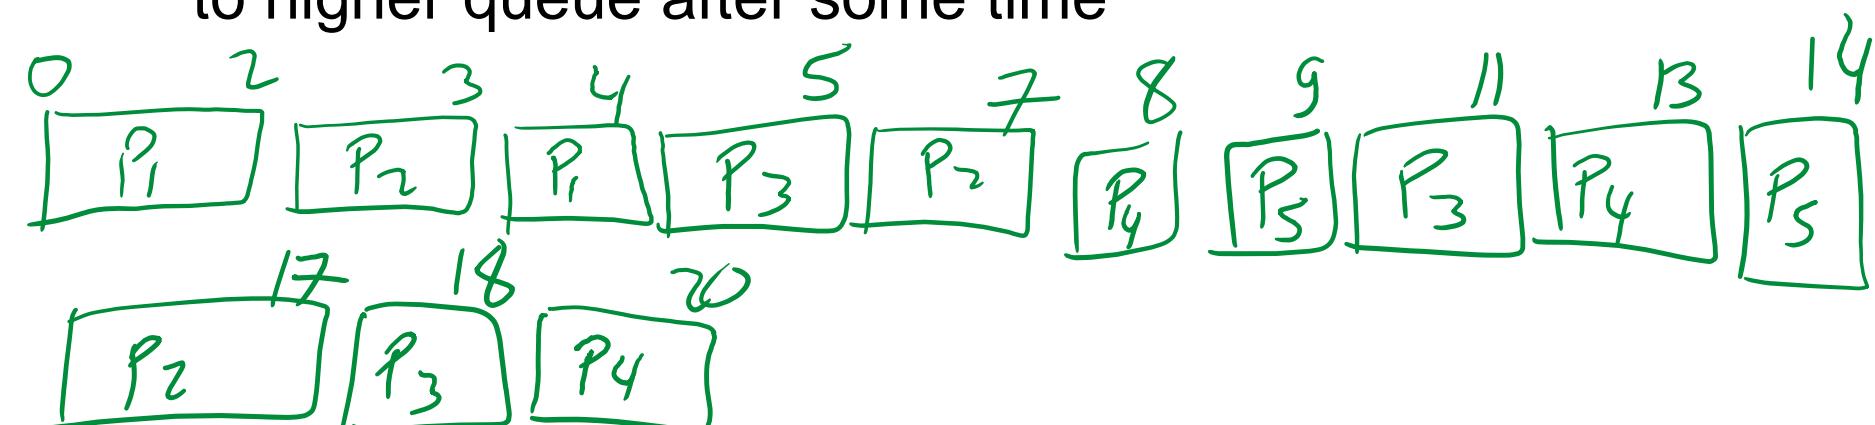
| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



Time Quantum for feedback Scheduling

- With a fixed quantum time, the turn-around time of longer processes can be high
- To alleviate this problem, the time quantum can be increased based on the depth of the queue
 - Time quantum of RQ_i = 2ⁱ⁻¹
- May still cause longer processes to suffer starvation.
 - Possible fix is to promote a process to higher queue after some time

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| 1 | 0 | 3 |
| 2 | 2 | 6 |
| 3 | 4 | 4 |
| 4 | 6 | 5 |
| 5 | 8 | 2 |



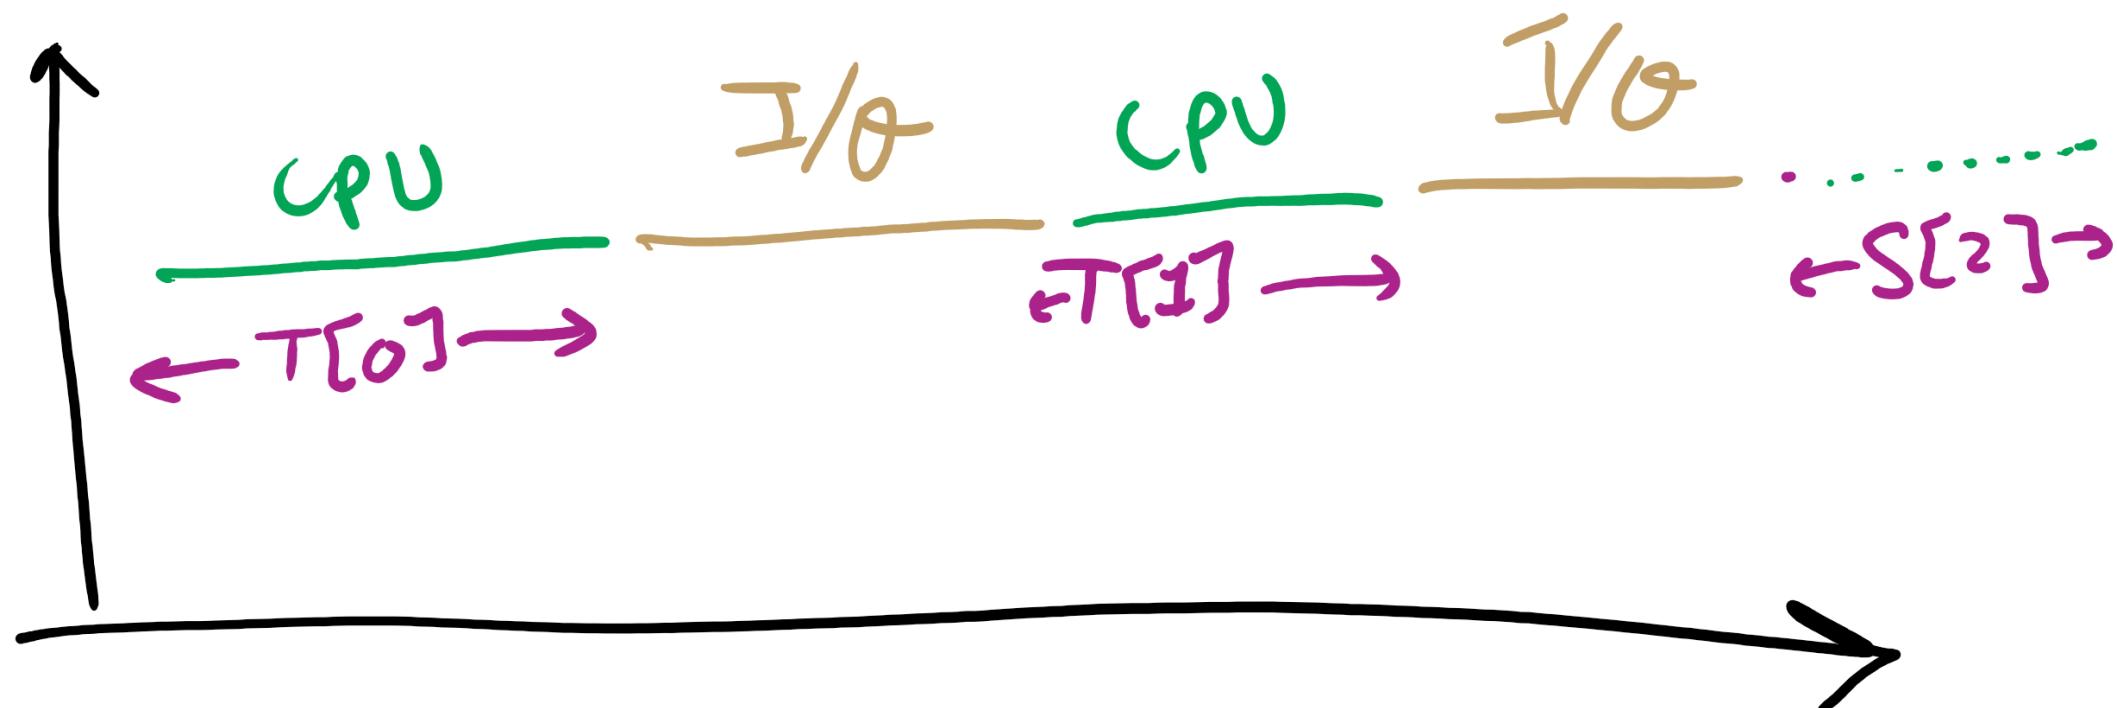
Algorithm Comparison

- Which one is the best?
- The answer depends on many factors:
 - the system workload (extremely variable)
 - hardware support for the dispatcher
 - relative importance of performance criteria (response time, CPU utilization, throughput...)
 - The evaluation method used (each has its limitations...)

Back to SJF: CPU Burst Estimation

- Let $T[i]$ be the execution time for the i th instance of this process: the actual duration of the i th CPU burst of this process
- Let $S[i]$ be the predicted value for the i th CPU burst of this process. The simplest choice is:
 - $S[n+1] = (1/n)(T[1] + \dots + T[n]) = (1/n) \sum_{\{i=1 \text{ to } n\}} T[i]$
- This can be more efficiently calculated as:
 - $S[n+1] = (1/n) T[n] + ((n-1)/n) S[n]$
- This estimate, however, results in equal weight for each instance

CPU Burst Estimation



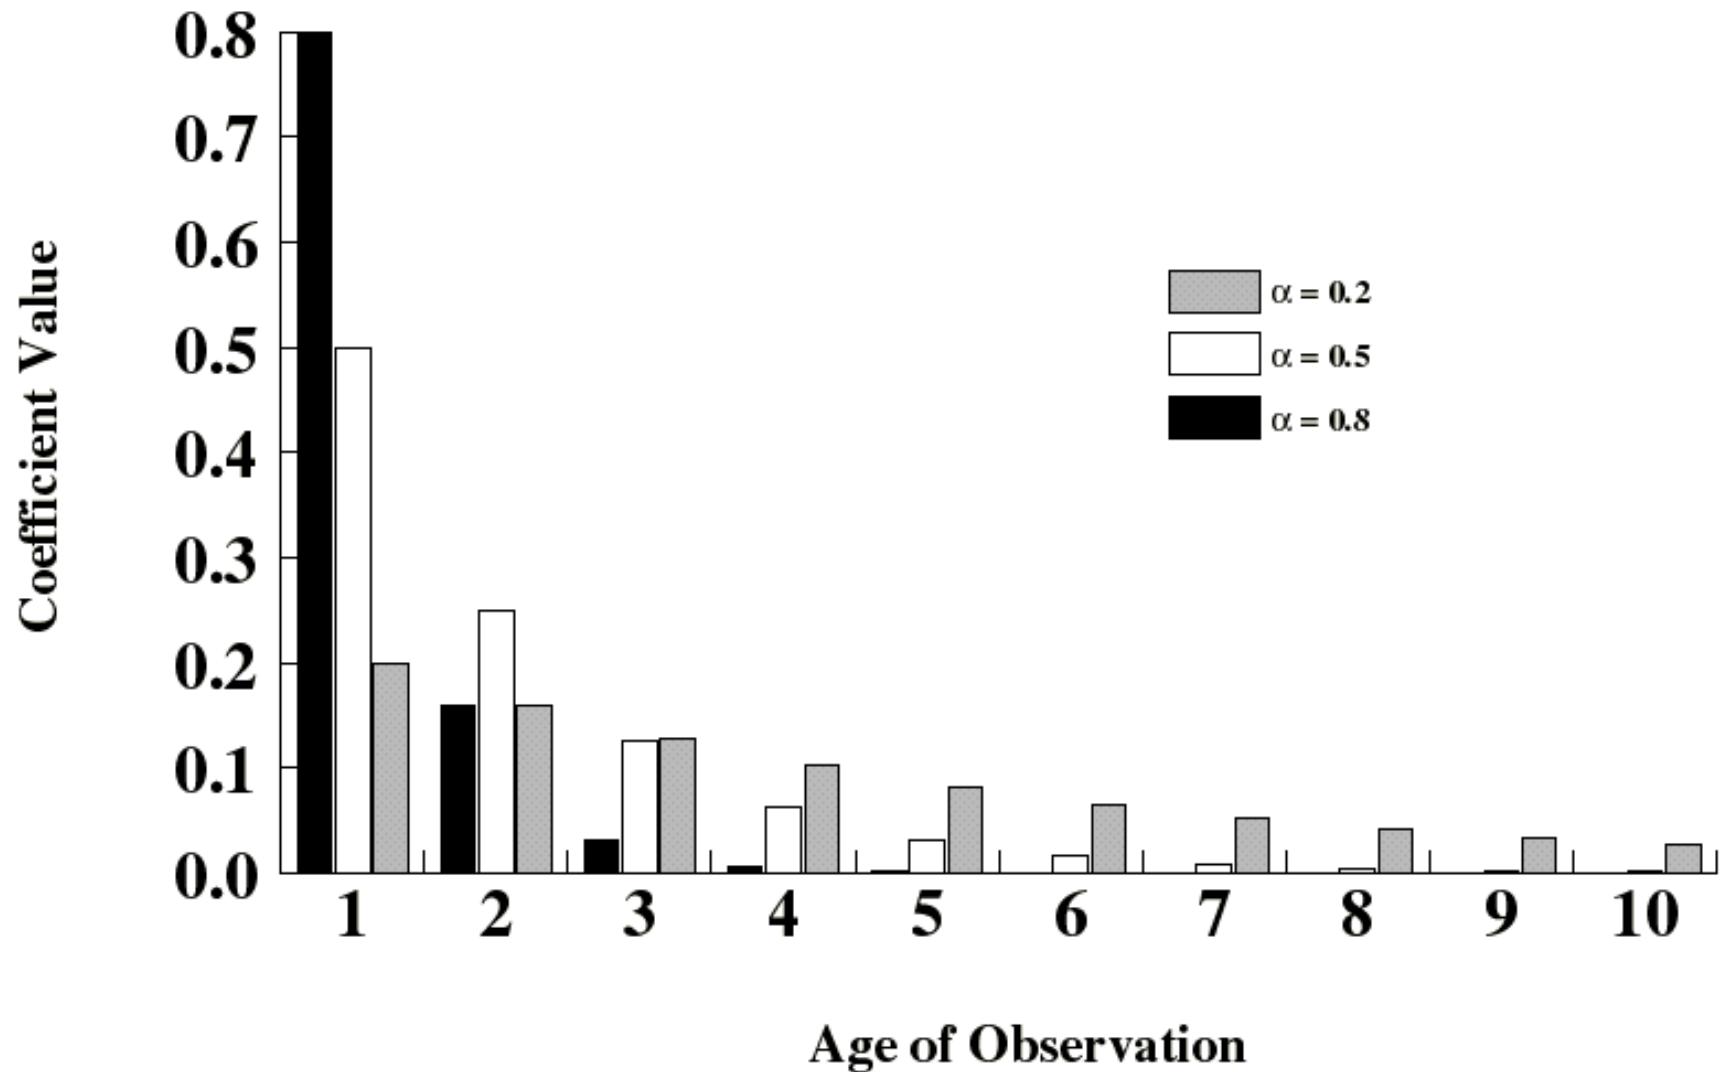
Estimating the required CPU burst

- Recent instances are more likely to better reflect future behavior
- A common technique to factor the above observation into the estimate is to use **exponential averaging** :
 - $S[n+1] = \alpha T[n] + (1 - \alpha) S[n]$; $0 < \alpha < 1$

CPU burst Estimate (Exponential Average)

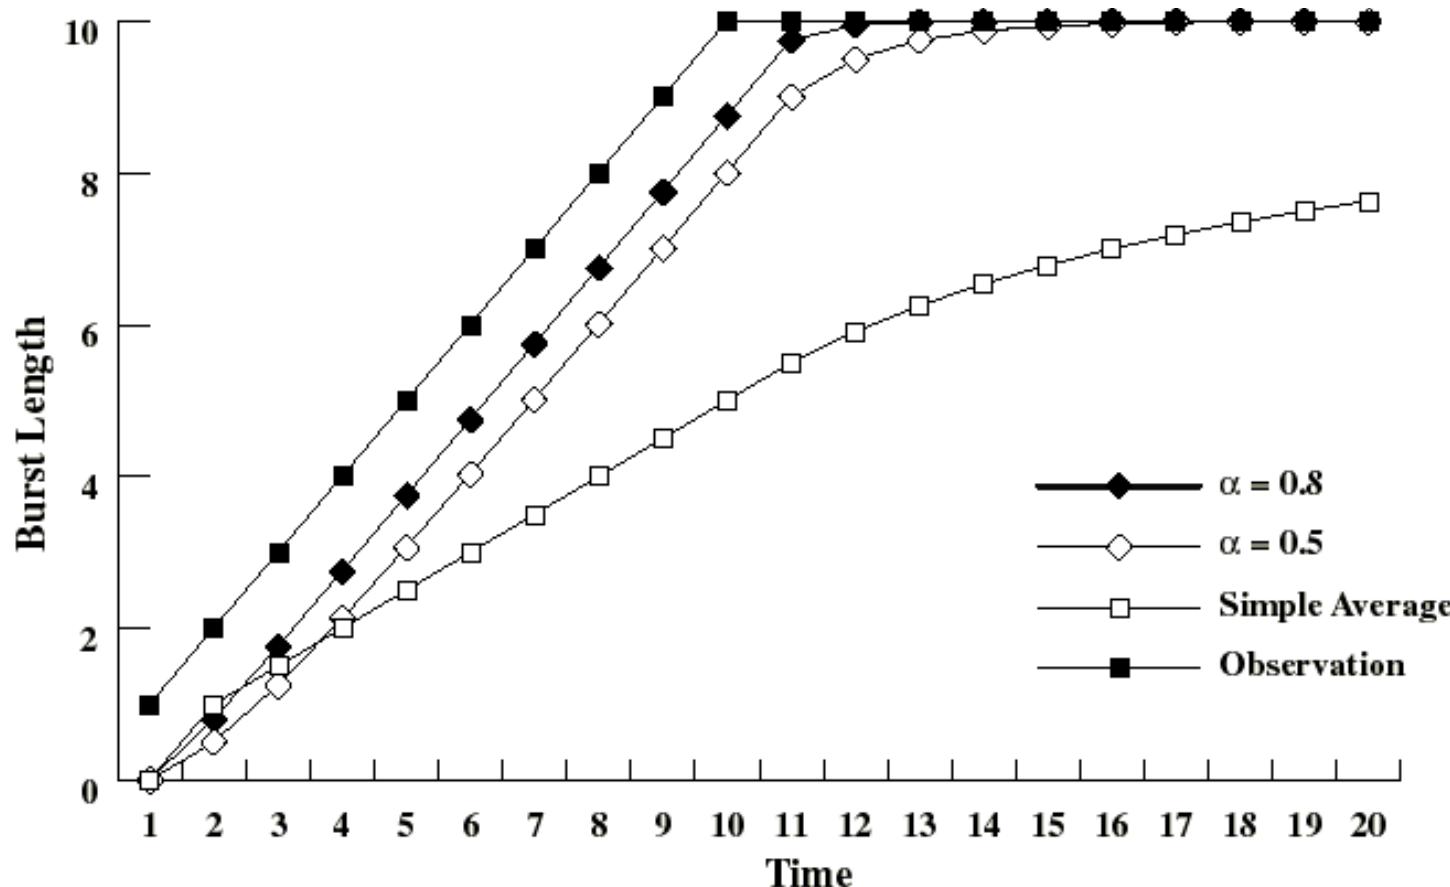
- Recent instances have higher weights, whenever $\alpha > 1/n$
- Expanding the estimated value shows that the weights of past instances decrease exponentially
 - $S[n+1] = \alpha T[n] + (1 - \alpha) \alpha T[n-1] + \dots (1 - \alpha)^{n-i} \alpha T[1] + \dots + (1 - \alpha)^n S[1]$
 - The predicted value of 1st instance, $S[1]$, is usually set to 0 to give priority to new processes

Exponentially Decreasing Coefficients

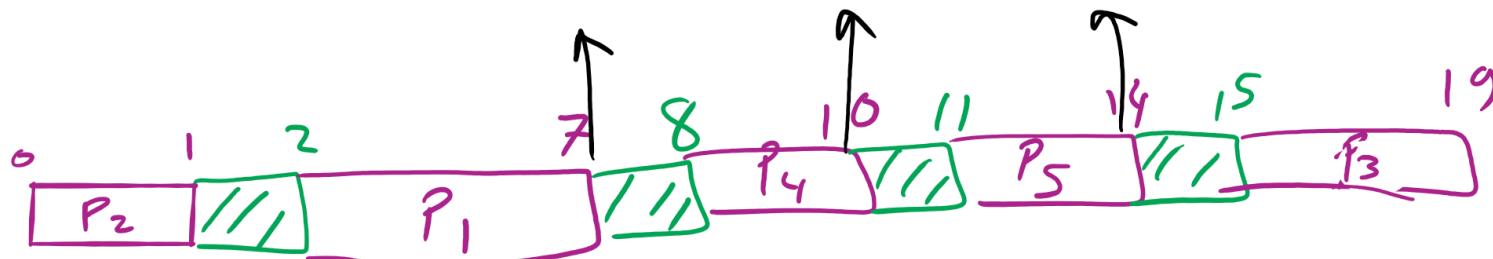


Exponentially Decreasing Coefficients

- $S[1] = 0$ to give high priority to new processes
- Exponential averaging tracks changes in process behavior much faster than simple averaging



FCFS Problem in HW7



$$\begin{aligned} \text{A.R.T.} &= \frac{P_1 + P_2 + P_3 + P_4 + P_5}{5} \\ &= \frac{5 + 1 + 11 + 6 + 7}{5} = 6 \text{ ms} \end{aligned}$$

$$\text{CPU Util} = \frac{19 - 4}{19} = \frac{15}{19}$$

CPU Burst Estimation

$$S[n+1] = \alpha T[n] + (1-\alpha) S[n]$$
$$0 < \alpha < 1$$

